

StateSet Agents

A Reinforcement Learning Framework
for Multi-Turn Conversational AI

Version 0.13.3 · May 2026

team@stateset.ai · github.com/stateset/stateset-agents

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

StateSet Agents: A Reinforcement Learning Framework for Multi-Turn Conversational AI

Technical Whitepaper

Version 0.13.3 · May 2026

StateSet Team · team@stateset.ai

Abstract

StateSet Agents is a reinforcement learning framework for training and serving large language model (LLM) agents that improve through **multi-turn interaction**. Unlike RLHF pipelines that optimize policies one response at a time, StateSet Agents treats the full conversational trajectory as the unit of optimization. The framework implements a family of **group-based policy optimization** algorithms — GRPO, GSPO, GEPO, DAPO, and VAPO — that reduce gradient variance by sampling multiple trajectories per prompt and computing advantages relative to a group baseline. This whitepaper describes the framework's architecture: the algorithmic foundations of each trainer, the agent and environment abstractions, the composable reward modeling system, and the operational layer (FastAPI serving, Helm/Kubernetes deployment, distributed training). We also present a comparative analysis of the five trainer variants and discuss the engineering tradeoffs that make conversational RL practical at scale.

⚠ Read this before `pip install`. The latest PyPI release is 0.7.1 (an early-2025 cut). It predates every named trainer in this paper, the Rust core, the dashboard, the auto-research loop, and most of the \$11.7 publication-gate machinery. Until a 0.13.x PyPI publish lands, install from source:

```
bash pip install git+https://github.com/stateset/stateset-agents@v0.13.3
```

If your tooling pins PyPI versions and you can't install from a tag, treat this paper as describing a *near-future* PyPI release rather than the current one. The framework's behavior is anchored to the named source commit, not to PyPI.

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

Versioning and Reproducibility

This whitepaper describes **version 0.13.3** of the framework. The implementation references — file paths, line numbers, default hyperparameters, LOC counts — are all taken from commit `4744c76` on `master`.

PyPI lag. At the time of writing, the latest PyPI release is **0.7.1**, which predates substantial parts of the surface described here (the named trainers, the Rust core, the dashboard, the auto-research loop). The 0.7.1 release also declares `Python >=3.8` in its classifiers, while the 0.13.3 source tree requires **Python ≥ 3.10** (with classifiers through 3.13) — when reading public PyPI metadata against this whitepaper, expect this gap. The full 0.13.3 surface can be obtained by installing from source (`pip install -e .` against the repository); a PyPI publication of 0.13.x is pending.

What's named here is anchored in code. Implementation citations (`gspo_trainer.py:390-419`, etc.) reference the named commit. To verify any specific claim:

```
git clone https://github.com/stateset/stateset-agents
cd stateset-agents
git checkout 4744c76          # the commit this whitepaper describes
grep -n "compute_sequence_importance_ratio" stateset_agents/training/gspo_trainer.py
```

Errata. Corrections published after this revision are tracked in [docs/WHITEPAPER_ERRATA.md](#). If `git log` shows commits more recent than `4744c76`, check the errata file before citing this document.

A complete reproducibility command list is in **Appendix C**.

First-party experimental results. §11.7 ("First-Party Reproduction") presents a canonical three-seed GSPO result on a customer-support task: Qwen2.5-0.5B-Instruct trained with the safe-default GSPOConfig produces +0.079 mean improvement on a paraphrase-tolerant LLM-judge with three-seed agreement (`benchmark_results/whitepaper_v1/customer_support_3seed_judge_qwen25_05b_instruct.json`). A negative-result companion artifact (Qwen3.5-0.8B trainee, baseline at ceiling 0.90) documents the framework's stability behavior when no headroom exists. Head-to-head trainer comparisons (GRPO vs GSPO vs DAPO vs VAPO) and vLLM speedup measurements remain pending and will land in a v1.1 revision.

How to Read This Document

This whitepaper is intentionally long ($\approx 10k$ words). Different audiences should navigate it differently:

- **Researchers** evaluating the algorithm coverage — §5 (Algorithmic Foundations) and §6.6 (Offline RL). The notation table in §5.0 and the comparative summary in §5.6 are the densest payoff.

Pinned to source commit. See Appendix C for reproducibility commands.

- **ML engineers** starting a training run → §11 (Practitioner's Guide), particularly the trainer-selection decision tree (§11.3) and the operational recipes table (§11.4). Then Appendix B for the hyperparameter defaults that ship with each trainer.
- **Platform engineers** considering deployment → §7 (Operational Layer) for endpoints, observability, dashboard, and Helm/K8s, plus §2 (System Architecture) for the layered design.
- **Framework maintainers** or contributors → §3 (Agent Abstractions), §4 (Reward Modeling), and §9 (Testing Philosophy) — the engineering invariants that keep the codebase coherent.
- **Anyone curious about the *design philosophy*** → §10 (Discussion) and §14 (Conclusion).

The appendices (file map, hyperparameter reference) are designed as standalone references — flip to them directly when you need a fact, not a narrative.

1. Introduction

1.0 Design Philosophy

Five principles recur throughout the framework. They are spelled out in detail with code-pattern references in §14, but it helps to state them up front because they constrain every later design decision:

1. **The trajectory is the unit of work.** Not the token, not the prompt — the trajectory. Every abstraction is shaped around this commitment.
2. **Algorithms are interchangeable; environments are not.** A `ConversationEnvironment` + reward function is a long-lived asset; choosing GSPO vs. DAPO is a tuning decision.
3. **Async-first, everywhere.** Rewards, environments, agent generation, tool calls — all async-native. This is what makes batched rollouts with composite rewards practical.
4. **Stubs are first-class.** `StubBackend` isn't a testing afterthought; it's the seam that makes the framework usable without GPUs.
5. **Failure surfaces are explicit.** Eight canonical exception tuples cover every retryable failure mode. Code that catches them retries; code that doesn't, doesn't.

If a section feels like it's pulling in a different direction from these, that's a bug — please file an issue.

Pinned to source commit. See Appendix C for reproducibility commands.

1.1 The Multi-Turn Problem

Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

Conventional Reinforcement Learning from Human Feedback (RLHF) — popularized by PPO-based training of InstructGPT-class models — treats each prompt-completion pair

as an independent episode. This works well for short, single-turn tasks, but it breaks down in two important regimes:

1. **Long chain-of-thought reasoning** (e.g., AIME-style math), where token-level importance ratios accumulate variance over thousands of tokens.
2. **Multi-turn dialogue and tool use**, where the policy's effect on episode reward is mediated by environment dynamics (user responses, tool outputs) and credit must be assigned across turns.

Recent research — notably GSPO (Qwen, 2025), DAPO (ByteDance, 2025), and VAPO (ByteDance, 2025) — has shown that **group-relative** updates dramatically improve stability. The core idea is to sample G trajectories per prompt, normalize rewards within each group (zero mean, unit variance), and use those normalized advantages as the policy gradient signal. This eliminates the need for a separate value network (in most variants) and reduces the variance that destabilizes token-level PPO on long sequences.

1.2 What StateSet Agents Provides

StateSet Agents packages this research into a coherent, deployable framework:

- **Agent abstractions** (`Agent` , `MultiTurnAgent` , `ToolAgent`) with pluggable Hugging Face / vLLM / stub backends.
- **Environments** (`ConversationEnvironment` , `TaskEnvironment`) that model multi-turn episodes with explicit `reset` / `step` semantics.
- **Trajectories** as first-class data structures carrying turn-level metadata, rewards, and tool calls.
- **Composable rewards** (`RewardFunction` , `CompositeReward`) supporting heuristic, domain, multi-objective, and neural scorers.
- **Group-based trainers**: GRPO, GSPO, GEPO, DAPO, VAPO — plus PPO and RLAIFF baselines.
- **Offline RL** (BCQ, BEAR, CQL, IQL, Decision Transformer) for learning from logged conversations.
- **Sim-to-real transfer, continual learning, long-term planning** modules.
- **Operational layer**: FastAPI service, OpenAI-compatible endpoints, Prometheus metrics, Helm charts, Kubernetes manifests.

Not all components are at the same level of production readiness. The framework ships with an explicit **Component Maturity matrix in §8.1** distinguishing *stable* (API-stable across point releases), *beta* (functionally complete, API may change), and *experimental* (works in tests, not yet recommended for production). New readers should check that matrix before designing a production deployment — VAPO, online GRPO, continual learning, and the auto-research loop are all marked experimental as of v0.13.3.

The framework is licensed under **BUSL-1.1**, converts to Apache 2.0 on 2029-09-03, distributed on PyPI as `stateset-agents` , and supports Python 3.10–3.13 on Linux and

Windows.¹ See [Component Maturity \(§8.1\)](#) for which surfaces are stable, beta, or experimental — relevant when planning a production deployment under this license.

2. System Architecture

2.1 Layered Design

[Mermaid diagram — for the rendered SVG see docs/WHITEPAPER.md on GitHub.]

```

flowchart TB
    API["FastAPI Service Layer<br>/v1/messages · /v1/chat/completions · /training · /metrics"]
    AgentSvc["Agent Service"]
    TrainSvc["Training Service"]
    InfSvc["Inference Service"]
    Mem["Memory<br>(Redis / SQL)"]
    BaseTrainer["BaseTrainer<br>GRPO · GSP0 · GEPO<br>DAPO · VAP0"]
    Agent["Agent<br>(MultiTurn / Tool)"]
    Env["Environment<br>+ Reward Function"]
    Backend["ModelBackend<br>(HF · vLLM · Stub)"]

    API --> AgentSvc
    API --> TrainSvc
    API --> InfSvc
    API --> Mem
    TrainSvc --> BaseTrainer
    AgentSvc --> Agent
    BaseTrainer --> Env
    InfSvc --> Backend
    Agent --> Backend
  
```

Each layer communicates through stable, typed interfaces. The serving and training layers can be deployed independently; the core abstractions (agents, environments, rewards) are usable as a Python library without any FastAPI or Kubernetes dependency.

2.2 Dependency Strategy

The package has a deliberately **lean core**: `numpy`, `pydantic`, `rich`, `typer`, `tqdm`, `cachetools`. Everything else — `torch`, `transformers`, `peft`, `trl`, `fastapi`, `vllm`, `optuna`, `deepspeed` — is an **optional extra**. This means:

- A user writing a custom reward function or environment never pays for a 2 GB torch install.
- CI smoke tests run on a minimal install; ML-heavy tests opt into `[training]`.
Pinned to source commit. See Appendix C for reproducibility commands.
- The stub backend (described in §3.3) makes the entire agent/environment stack runnable without any model weights.
Licensed B2 SL 1.1 (transitioning to Apache 2.0 on 2025-01-01).

Extras are: `[training]`, `[api]`, `[trl]`, `[vllm]`, `[hpo]`, `[distributed]`, `[examples]`, `[auto-research]`.

3. Agent and Environment Abstractions

The class relationships in this section are dense — here is a single map of how the pieces connect before we describe each:

[Mermaid diagram — for the rendered SVG see docs/WHITEPAPER.md on GitHub.]

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

```

classDiagram
class Agent {
    +AgentConfig config
    +ModelBackend backend
    +initialize() async
    +generate_response(messages, context) async
}
class MultiTurnAgent {
    +ConversationContext context
    +PlanningManager planning
    +ConversationMemory memory
    +reset() async
    +register_tool(tool)
}
class ToolAgent {
    +list tools
    +dict tool_registry
    +add_tool(tool)
}
class ModelBackend {
    <<Protocol>>
    +tokenizer
    +model
    +generation_config
}
class HuggingFaceBackend
class VLLMBackend
class StubBackend

class EnvironmentBase {
    +reset(scenario) async
    +step(state, action) async
    +run_episode(agent, scenario, max_turns) async
}
class ConversationEnvironment {
    +list scenarios
    +RewardFunction reward_fn
    +clone()
}

class RewardFunction {
    +compute_reward(turns, context) async
    +weight: float
}
class CompositeReward {
    +list components
    +str combination
}

class ConversationTurn {
    +role
    +content
    +reward
    +tool_calls
    +metadata
}
class MultiTurnTrajectory {
    +list turns
    +episode_id

```

Pinned to source commit. See Appendix C for reproducibility commands.
 Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).


```

    +total_reward
  }

  Agent <|-- MultiTurnAgent
  MultiTurnAgent <|-- ToolAgent
  Agent o-- ModelBackend
  ModelBackend <|.. HuggingFaceBackend
  ModelBackend <|.. VLLMBackend
  ModelBackend <|.. StubBackend
  EnvironmentBase <|-- ConversationEnvironment
  ConversationEnvironment o-- RewardFunction
  RewardFunction <|-- CompositeReward
  CompositeReward o-- RewardFunction : aggregates
  MultiTurnTrajectory o-- ConversationTurn

```

The trainer (§5) sits above this diagram: it drives the environment, samples the agent, scores the resulting trajectories via the reward function, and updates the model wrapped inside the backend.

3.1 The Agent Hierarchy

The framework defines a three-tier agent hierarchy in `stateset_agents/core/`:

Class	File	Role
<code>Agent</code>	<code>core/agent.py</code>	Base class. Owns a <code>ModelBackend</code> , handles lazy initialization, exposes <code>generate_response</code> .
<code>MultiTurnAgent</code>	<code>core/multiturn_agent.py</code>	Adds <code>ConversationContext</code> , optional <code>DialogueDatabase</code> , planning manager, and <code>async reset()</code> .
<code>ToolAgent</code>	<code>core/tool_agent.py</code>	Adds OpenAI-compatible function calling, <code>tool_registry</code> , JSON schema validation, parallel tool execution.

All response generation follows a single async signature:

```

async def generate_response(
    messages: str | list[dict[str, str]],
    context: dict[str, Any] | None = None,
) -> str

```

This uniformity allows the same training loop to drive plain chat agents, tool-using agents, and planning agents without branching.

Tool-call semantics in training. When `ToolAgent` makes a tool call during a rollout, four pieces of behavior need to be pinned down because they materially affect what the trainer optimizes: Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

- **A tool call is *part of* an assistant turn, not a separate turn.** The `ConversationTurn` for the assistant carries `tool_calls: list[ToolCall]` and `tool_results: list[ToolResult]` as fields alongside `content`. The trainer's response-

token mask covers the assistant text *and* the JSON-encoded tool call; the tool *result* is masked out (it's a "tool" role from the environment, not from the policy).

- **Credit assignment is per-turn, not per-call.** Group-relative advantages are computed at the turn level. A turn that contains "say X" + "call tool Y with args Z" gets one scalar advantage that flows through all the policy-emitted tokens in that turn. There is no per-tool-call advantage signal; cross-tool credit assignment is an open problem (§10.6).
- **Tool errors are graded by the reward function, not the trainer.** The framework's design choice: if a tool raises, the `ToolResult.error` field is populated; the trainer makes no policy-level distinction between "successful call" and "errored call." Whether that's penalized is the reward function's job — typically via a component like `TaskCompletionReward` or a custom rubric that reads `tool_results`.
- **Tool outputs are excluded from the policy gradient.** Only assistant-emitted tokens contribute to the loss. Tool outputs (potentially long, often verbatim API responses) sit in the context for the next turn but aren't optimized over. This is what allows training against tools with large or non-stationary output spaces (search results, database query responses) without polluting the gradient.

These semantics are stable in v0.13 and apply to all group-based trainers. Tool-use as a first-class RL signal — per-call advantages, tool-selection learning under uncertainty — is the open work in §10.6.

3.2 Configuration: `AgentConfig`

`AgentConfig` (defined in `core/agent_config.py`) is a 22-field dataclass covering model selection, generation hyperparameters, PEFT/LoRA settings, reasoning tags (DeepSeek-R1-style), and planning configuration. Validation logic enforces bounds and emits remediation suggestions for misconfigurations.

Two fields are critical for the framework's testability:

- `use_stub_model: bool = False`
- `stub_responses: list[str] | None = None`

Setting `use_stub_model=True` skips all Hugging Face loading and substitutes a deterministic in-memory backend. This is what enables the 2,438-test suite to run in seconds without GPU hardware.

3.3 Backends: The `ModelBackend` Protocol

`ModelBackend` is a runtime-checkable Protocol (PEP 544) declaring three properties — `tokenizer`, `model`, `generation_config` — each backed by its own protocol. The framework ships three concrete backends:

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

1. **Hugging Face backend** — loads `AutoModelForCausalLM` + `AutoTokenizer` with configurable dtype, attention implementation (`flash_attention_2`, `sdpa`, `eager`), and device map. Supports 4-bit/8-bit quantization via `bitsandbytes` and LoRA adapters via `peft`.

2. **vLLM backend** — uses `vllm.LLM` for batched, paged-attention generation [14]. Typically delivers a large speedup over plain HF `model.generate` during rollout collection; magnitude depends heavily on model, batch size, and sequence length (see §6.4 for measurement guidance). Surfaces token log-probs directly from the vLLM sampler so the trainer can skip a redundant forward pass.
3. **Stub backend** — `StubModel` + `StubTokenizer` + `StubGenerationConfig`. Returns canned or templated responses; tokenizes with a fixed 256-character vocabulary. No external downloads, no GPU.

Backend selection happens in `Agent.initialize()`: an injected backend wins, otherwise `use_stub_model=True` selects the stub, otherwise the framework loads from Hugging Face. This **dependency injection** pattern (over global patching) is what keeps the test suite reliable as the codebase grows.

3.4 Environments

`EnvironmentBase` (in `core/environment_base.py`) defines the canonical RL interface:

```
async def reset(scenario: dict | None) -> EnvironmentState
async def step(state: EnvironmentState, action: ConversationTurn)
    -> tuple[EnvironmentState, float, bool, dict]
async def run_episode(agent_fn, scenario, max_turns)
```

`EnvironmentState` tracks `episode_id`, `turn_count`, an `EpisodeStatus` enum (ONGOING / COMPLETED / FAILED / TIMEOUT), and a `context` dict.

`ConversationEnvironment` extends the base with scenario-driven dialogue: each `reset` selects a scenario (target user persona, conversation goal, evaluation criteria), and each `step` advances the conversation by one agent turn, generates a user response, computes the step reward, and decides whether to terminate. A `clone()` factory supports parallel environment workers.

3.5 Trajectories and Turns

The atomic unit of data is `ConversationTurn`, supporting both a legacy `(user_message, assistant_response, reward)` tuple and a modern `(role, content, reward, tool_calls, tool_results)` form. Turns aggregate into `MultiTurnTrajectory` objects that carry episode metadata, total reward, and the full ordered turn list. Trainers consume trajectories; rewards return `RewardResult` objects that decompose into a scalar `score` plus a `breakdown` dict.

3.6 The Memory Subsystem

Pinned to source commit. See Appendix C for reproducibility commands.

Skim on first read — relevant only for production deployments where conversations exceed the model's native context window. Algorithm details matter when you're integrating Redis / SQLite backends; the trainer treats memory as transparent context.

Long conversations exceed the context windows of even modern models. `core/memory.py` provides a multi-tier memory system that lets agents reference earlier context selectively, without naively passing every prior turn into every generation call.

`MemoryConfig` declares the tiers:

- **SHORT_TERM** — recent turns within a sliding window (bounded by `max_short_term_turns` and `max_short_term_tokens`). Always included verbatim in the next prompt.
- **LONG_TERM** — persistent facts and summaries surviving across episodes. Retrieved by relevance.
- **EPISODIC** — per-episode summaries; the unit of consolidation when a conversation closes.
- **SEMANTIC** — extracted entities and facts (names, dates, IDs, preferences) for structured lookup.
- **WORKING** — current-task context: open subgoals, intermediate results, pending tool calls.

The actual algorithm — `ConversationMemory.add_turn` in `core/memory.py` executes the following sequence on every turn:

1. **Append** the turn to short-term with a per-turn importance score (default 0.5) and ISO timestamp.
2. **Extract entities** (default: regex over name / date / id / email / phone patterns) into a `dict[entity_key → Entity]`. Duplicates are deduplicated by `f"{entity_type}:{value.lower()}"`.
3. **Extract facts** (default: heuristic sentence-level filter for declarative statements). Facts are deduplicated and capped at the most recent 50.
4. **Summarize** if `turn_count - last_summarization >= summary_threshold` (default 10): an async `_maybe_summarize()` is scheduled in the background, which compacts the oldest `summary_threshold` turns into a single `Summary` object and pushes it into the long-term tier.
5. **Decay** the `importance` of every turn already in short-term by `importance_decay` (default 0.95) — exponential demotion over time.
6. **Trim** short-term to at most `max_short_term_turns` (default by token budget); turns trimmed with `importance > 0.3` get promoted to long-term, the rest are discarded.

Retrieval (`get_context_for_generation`) reverses through short-term newest-first, accumulating turns until `max_short_term_tokens` is reached, then optionally appends the latest summary and the entity/fact tables. Long-term semantic retrieval (top-k by relevance) is supported via `get_relevant_memories(query)` but is *not* automatically wired into `get_context_for_generation` — callers opt in.

Persistence backends are pluggable: in-memory (default, ephemeral), Redis (for distributed agent fleets), or SQLite (for single-host durability). Long-term memory survives process restarts under Redis/SQLite, which is what enables continuity across sessions for production deployments. **Pending engineering work:** the semantic-tier

Pinpoint to source commit. See Appendix C for reproducibility commands.
Licensed BY/SL-1.1 (transitioning to Apache 2.0 on 2029-09-03)

retrieval uses string-overlap heuristics by default rather than embedding similarity; swapping in a sentence-transformer-backed semantic store is the v1.1 follow-up. Until that lands, treat semantic retrieval as "useful for entity/fact pulls" but not "useful for paraphrase retrieval of arbitrary content."

4. Reward Modeling

4.1 The Reward Function Hierarchy

All rewards inherit from `RewardFunction` and implement:

```
async def compute_reward(
    turns: list[ConversationTurn],
    context: dict[str, Any] | None = None,
) -> RewardResult
```

`RewardResult` is the canonical return type:

```
@dataclass
class RewardResult:
    score: float
    breakdown: dict[str, float]
    metadata: dict[str, Any]
    explanation: str | None
```

The `total_reward` property aliases `score` for backwards compatibility with older trainer code. Both forms serialize losslessly via `to_dict()` / `from_dict()`.

4.2 Composition

`CompositeReward` combines multiple reward functions via four aggregation methods — `weighted_sum` (default), `average`, `min`, `max` — with graceful degradation: if any component raises, the composite continues with the remaining components and logs the failure rather than aborting training.

The framework ships a battery of concrete reward functions in `core/basic_rewards.py` and `core/domain_rewards.py`:

Reward	Signal
<code>HelpfulnessReward</code>	LLM-judge helpfulness rating
<code>SafetyReward</code>	Pinned to source commit. See Appendix C for reproducibility commands. Licensed BSL 1.1 (transitioning to Apache 2.0 on 2029-09-03). Toxicity, harmful content filter
<code>CorrectnessReward</code>	Verifier match against ground truth
<code>ConcisenessReward</code>	Length penalty, redundancy detection

Reward	Signal
EngagementReward	Heuristic engagement score
TaskCompletionReward	Goal-state predicate satisfaction
CustomerServiceReward	Domain composite (resolution + tone + accuracy)
TechnicalSupportReward	Domain composite (correctness + step-by-step structure)
SalesAssistantReward	Domain composite (qualification + conversion intent)

4.3 Multi-Objective Rewards

`MultiObjectiveReward` (in `rewards/multi_objective_reward.py`) supports per-turn and per-trajectory components evaluated in parallel via `asyncio.gather`. This is the standard pattern for production deployments where the scalar reward is a weighted sum of correctness, safety, brand voice, and latency penalties.

4.4 Neural Reward Models

Hand-crafted rewards have a ceiling: they capture exactly what their designer encoded. When you have labeled examples — pairwise preferences, scored conversations, RLHF demonstrations — you can do better by *learning* the reward function. `training/neural_reward_trainer.py` (687 LOC) provides this pathway.

The `NeuralRewardModel` (lines 137-176) is a deliberately small MLP:

- **Input.** Concatenated prompt + response embeddings (128-dimensional by default). The default `encoder=None` path uses a hash-based pseudo-embedding that exists **purely to smoke-test the training loop**. It cannot learn a useful reward function: the same string always hashes to the same vector and that vector is uncorrelated with text semantics, so the model has nothing to fit against. The constructor emits a `WARNING` when `encoder=None` for exactly this reason. **For real reward learning**, pass `encoder=` a `sentence_transformers.SentenceTransformer` or any callable `text -> tensor(embedding_dim,)`. The framework will not catch you if you ignore the warning, but the resulting loss curve will look like noise — the failure is loud in W&B but silent in the absence of monitoring. Make `sentence-transformers` part of your requirements if you use this path.
- **Body.** A configurable stack of hidden layers (default 2 layers, `hidden_dim=256`) with ReLU/GELU activations and `dropout=0.1`.
- **Output.** A single scalar (the reward).
- **Initialization.** Xavier uniform for stable early training.

Pinned to source commit. See Appendix C for reproducibility commands.

The trainer fits this network to a labeling source — typically the output of an existing `CompositeReward` running offline, or scored data from a stronger LLM-judge — and produces a fast, differentiable, GPU-resident reward function that doesn't require network calls during rollout scoring. This is the "Bitter Lesson" [Sutton, 2019][20] applied to rewards: stop encoding domain knowledge by hand once you have enough

data to learn it. The caveat is that *learned* rewards have well-known failure modes — reward hacking against the learner's blind spots, judge gaming, and miscalibration on distribution shift — that hand-crafted composites are structurally resistant to. The framework's bias toward composites as the *default* (with learned rewards as the escalation when data scale justifies it) is a deliberate hedge, not a stylistic choice.

Use cases: speeding up training when an LLM-judge reward is the bottleneck (a hosted gpt-4o-class judge can dominate step time at >\$50/run); enabling differentiable reward shaping in algorithms that want gradient flow through the reward signal.

4.5 RLAIF and Constitutional Critique

`training/rlaif_trainer.py` (749 LOC) implements RLAIF [10] augmented with a Constitutional AI critique-revision loop [9]. The `ConstitutionalAI` class (lines 170-198) implements the three-step inner loop on every training sample:

1. **Generate** an initial response from the current policy.
2. **Critique** the response against a domain-specific set of principles. The framework ships four principle sets — `general`, `assistant`, `coding`, `reasoning` — in `CONSTITUTIONAL_PRINCIPLES` (lines 146-167); custom principle sets can be plugged in via config.
3. **Revise** the response in light of the critique.

The revised response is then scored by a judge model. `RLAIFConfig` selects the judge: `judge_provider` $\in$ `{openai, anthropic, local}`, `judge_model` defaults to `gpt-4o`. Scoring criteria are `helpfulness`, `correctness`, `harmlessness`.

The resulting scalar reward feeds into a downstream policy optimizer — RLAIF can be configured to use PPO, GRPO, or GSPO under the hood (`config.algorithm`, line 118), with sensible defaults: `beta=0.1` (KL penalty enabled — critical when the judge is more powerful than the learner), `clip_eps=0.2`, `num_generations=4`, optional self-play every 100 steps.

This is the canonical path for training agents without human-labeled preference data: bootstrap from a stronger model's judgment, regularize against drift via KL to a reference, and iterate.

5. Algorithmic Foundations

This section describes the five group-based policy optimization algorithms that form the core of the training stack. All inherit from `BaseTrainer` (in `training/base_trainer.py`), which provides shared optimizer setup (AdamW with cosine warmup), reference-model KL computation, gradient clipping with statistics tracking, mixed-precision support, and W&B integration.

5.0 Notation

Symbol	Meaning
x	Prompt (the conditioning context)
$y_i = (y_1, \dots, y_{L_i})$	The i -th sampled response in a group
G	Group size: number of responses sampled per prompt
$L_i = y_i $	Token length of response i
π_θ	Current (learner) policy with parameters θ
π	Behavior (sampler) policy — the policy used to draw rollouts
π_{ref}	Frozen reference policy, typically the SFT initialization
r_i	Scalar reward for trajectory i , returned by the reward function
A_i	Group-relative advantage: $(r_i - \mu) / \sigma$ where μ, σ are the group's mean and std
ρ_t	Token-level importance ratio $\pi_\theta(y_t \mid \cdot) / \pi(y_t \mid \cdot)$
s_t	Sequence-level importance ratio (GSPO): $\left(\prod_{\tau=t}^T \rho_\tau \right)$
$\epsilon, \epsilon_L, \epsilon_H$	PPO clip parameters; $\epsilon_L \neq \epsilon_H$ under Clip-Higher
β	KL penalty coefficient (often 0 in group-based methods)
$V_\phi(s_t)$	Value function (only in VAPO)
λ	GAE decay parameter (only in VAPO)

All trainers operate over **response tokens only**: prompt tokens are masked out of every sum, log-prob, and loss term.

5.1 GRPO (Group Relative Policy Optimization)

Source. The original group-based RL formulation, popularized by DeepSeek-Math [5]. StateSet Agents wraps the Hugging Face TRL implementation via `TRLGRPOTrainerWrapper`.

Objective. For each prompt x , sample G responses $\{y_i\}_{i=1}^G$. Compute group-normalized advantages $A_i = (r_i - \mu) / \sigma$, then optimize the standard clipped surrogate at the token level:

$$\mathcal{L}_{\text{GRPO}} = -\mathbb{E} \left[\min \left(\rho_t A_i, \text{clip}(\rho_t, 1 - \epsilon, 1 + \epsilon) A_i \right) \right]$$

$1 - \frac{1}{1 + \epsilon} A_i$

where $\rho_t = \pi_\theta(y_t \mid x, y_{<t}) / \pi_{\theta_{\text{old}}}(y_t \mid \cdot)$.

Defaults. `num_generations=4`, `beta=0.0` (KL disabled), `max_grad_norm=1.0`, `gradient_checkpointing=True`.

When to use. General-purpose baseline. Best when the reward verifier is reliable and trajectories are short-to-medium length.

5.2 GSPO (Group Sequence Policy Optimization)

Source. Zheng et al. [1], Qwen team. StateSet Agents implements this in-house: `GSPOTrainer` (`training/gspo_trainer.py`, 852 LOC).

Key innovation. Replace the token-level importance ratio with a **length-normalized sequence-level** one:

$$s_i(\theta) = \left(\frac{\pi_\theta(y_i \mid x)}{\pi_{\theta_{\text{old}}}(y_i \mid x)} \right)^{1/|y_i|} = \exp \left(\frac{1}{|y_i|} \sum_t \log \frac{\pi_\theta(y_{i,t} \mid \cdot)}{\pi_{\theta_{\text{old}}}(y_{i,t} \mid \cdot)} \right)$$

The objective then applies standard symmetric PPO clipping over the sequence-level ratio:

$$\mathcal{L}_{\text{GSPO}} = -\mathbb{E} \left[\min \left(s_i(\theta) A_i, \frac{A_i}{s_i(\theta)}, 1 - \epsilon_L, 1 + \epsilon_H \right) A_i \right]$$

The framework allows $\epsilon_L \neq \epsilon_H$ in the config, but the shipped defaults are symmetric — Clip-Higher asymmetry is reserved for DAPO and VAPO (§5.4–5.5). The critical thing about the GSPO clip bounds is not their symmetry but their *magnitude*: because s_i is already exp-of-a-small-per-token-quantity, the bounds must be much tighter than token-level PPO's 0.2. See the defaults note below.

Why it matters. Token-level ratios accumulate variance multiplicatively across the sequence; on long outputs and Mixture-of-Experts models, this manifests as training collapse. Length normalization keeps the importance weight in a stable regime regardless of $|y_i|$.

Pseudocode (paraphrased from `gspo_trainer.py:390-680`):

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

```
# Phase 1: rollout collection
for prompt in batch:
    responses = generate(policy_old, prompt, n=group_size)
    rewards = await reward_fn(responses)
    advantages = (rewards - rewards.mean()) / (rewards.std() + 1e-8)

# Phase 2: GSPO update
for response, A in zip(responses, advantages):
    L = response_length(response)
    logp_new = sum_logprob(policy_new, response) # gradient-tracked
    logp_old = sum_logprob(policy_old, response) # detached
    s = torch.exp((logp_new - logp_old) / L) # sequence-level ratio
    obj = torch.min(s * A, s.clamp(1-eps_L, 1+eps_R) * A)
    loss = -obj.mean()
    if beta > 0 and ref_model is not None:
        loss = loss + beta * forward_kl(policy_new, ref_model, response) / L
    loss.backward(); optimizer.step()
```

Implementation citations. Sequence ratio computed at `gsपो_trainer.py:390-419` (log-sum, length-normalize, exponentiate as separate steps to retain numerical stability). Clipped surrogate at `gsपो_trainer.py:639-649`. Per-sequence KL penalty (only when `beta > 0`) at lines 652-660. Right-padding is enforced for stable prompt-boundary detection across the batch.

Defaults. `num_generations=4`, `clip_range_left=3e-4`, `clip_range_right=4e-4`, `warmup_ratio=0.1`. These clip bounds are taken from the original GSPO paper and are roughly three orders of magnitude tighter than token-level PPO — necessary because the length-normalized ratio lives close to 1.0. **If you see no exploration, this is the first knob to widen.** Single gradient step per rollout (no inner PPO epochs) — a deliberate choice for stability with on-policy data.

When to use. Long outputs (CoT, code, structured generation), MoE models, any case where token-level GRPO is unstable.

5.3 GEPO (Group Expectation Policy Optimization)

Source. Reference [4]. Implementation in `training/gepo_trainer.py` (724 LOC).

Key innovation. Replace per-trajectory importance weights with a **group-level**

expectation. Let $p(y \mid x)$ denote the learner-policy sequence probability and $q_i = q(y_i \mid x)$ the sampler-policy probability for the i -th group member. Define:

$$w_{\text{GEIW}}(y \mid x) = \frac{p(y \mid x)}{E_q[q]}, \quad E_q[q] \approx \sum_{i=1}^G \tilde{q}_i \cdot q_i, \quad \tilde{q}_i = \frac{q_i}{\sum_{i=1}^G q_i}$$

Pinned to source commit. See Appendix C for reproducibility commands.

The denominator $E_q[q]$ is the **self-normalized importance sampling**

denominator under the empirical normalized distribution $\{\tilde{q}_i\}$ over the group — same primitive that appears in particle-filter resampling and population Monte Carlo. It is computed once per group as a scalar; every member of the group is divided

by the same value, which is what amortizes importance-weight variance across the group. Standard PPO clipping is then applied over w .

Why it matters. Group-level aggregation amortizes importance-weight variance across the group, which is especially valuable when sampler and learner policies diverge (e.g., heterogeneous compute, network-delayed actors, off-policy data). The denominator is scalar, avoiding per-sequence division instabilities.

Implementation citations. `GEPOTrainer` in `gepo_trainer.py` (724 LOC). The group-expectation denominator is computed at `gepo_trainer.py:301-337` (`compute_gepo_coefficient`); the sampler-side probabilities are explicitly **detached** at line 326 so the denominator does not flow gradient. The computation is per-sequence (1-D tensors of shape `[group_size]`), so every member of a group divides by the *same* scalar expectation — this is what amortizes variance across the group. Clipped surrogate at lines 517-534.

Defaults. `group_size=8`, `clip_eps=0.2`, `beta=0.0` (KL typically disabled — group weights handle divergence implicitly), `use_group_baseline=True`.

When to use. Distributed or asynchronous training, off-policy batches, replay-based RL.

5.4 DAPO (Decoupled Clip and Dynamic Sampling Policy Optimization)

Source. Yu et al. [2], ByteDance. Implementation in `training/dapo_trainer.py` (945 LOC). The paper reports 50/60 on AIME 2024 with Qwen-2.5-32B; we have not independently reproduced this number.

Four innovations.

1. **Clip-Higher (asymmetric clipping).** $\text{clip}(\rho_t) \in [1 - \epsilon_L, 1 + \epsilon_H]$ with $\epsilon_L = 0.2, \epsilon_H = 0.28$. Allowing more upside than downside encourages exploration without sacrificing stability.
2. **Token-level loss normalization.** Divide the surrogate sum by **total response tokens** across the batch, not by sample count. This prevents the implicit bias toward shorter sequences that arises when each sample contributes equally regardless of length:

\$\$

$$\mathcal{L}_{\text{DAPO}} = -\frac{1}{\sum_i |y_i|} \sum_i \min(\rho_{i,t} A_i, \text{clip}(\rho_{i,t}) A_i)$$

\$\$

1. **Dynamic sampling.** A `DynamicSamplingBuffer` filters out prompts where all G rollouts are correct or all are incorrect — these provide zero gradient signal. The buffer keeps sampling until every retained prompt has $0 < \text{accuracy} < 1$.
2. **Overlong reward shaping.** A graduated length penalty: no penalty if $|y| \leq L_{\text{max}} - L_{\text{cache}}$; linear interpolation from 0 to -1 within the cache region; full -1 penalty if $|y| > L_{\text{max}}$. Defaults: `max_generation_length=20480`, `overlong_cache_length=4096`.

Implementation citations. DAPOTrainer in dapo_trainer.py (945 LOC).

DynamicSamplingBuffer.should_include at lines 276-278 (strict $\text{min_accuracy} < \text{acc} < \text{max_accuracy}$; both 0.0 and 1.0 are excluded). DAPORewardShaper.compute_length_reward at lines 226-244 — three-region piecewise penalty with $\text{soft_start} = \text{max_length} - \text{cache_length} = 16384$ by default. Asymmetric clip and token-level loss normalization at lines 542-567. Optional inner gradient updates loop at line 850 with `old_token_log_probs.detach()` at line 858.

Pseudocode (paraphrased):

```
# Phase 1: dynamic sampling – keep rolling until every prompt has  $0 < \text{acc} < 1$ 
buffer = DynamicSamplingBuffer(min_accuracy=0.0, max_accuracy=1.0)
while not buffer.full(batch_size):
    prompt = sample_prompt()
    responses = generate(policy_old, prompt, n=group_size)
    rewards = await reward_fn(responses)
    accuracy = (rewards > threshold).float().mean()
    if  $0.0 < \text{accuracy} < 1.0$ :
        # Apply overlong shaping per response
        for r, response in zip(rewards, responses):
            L = len(response)
            if L > max_length:
                r += -1.0 # full penalty
            elif L > soft_start:
                r += -(L - soft_start) / cache_length # linear penalty
        buffer.add(prompt, responses, rewards)

# Phase 2: DAPOTrainer update – token-level, Clip-Higher
total_tokens = sum(response_lengths(buffer))
for response, A in buffer:
    for t in response_tokens(response):
        r_t = exp(logp_new[t] - logp_old[t])
        obj_t = min(r_t * A, r_t.clamp(1-eps_L, 1+eps_H) * A)
        loss = -sum(obj_t) / total_tokens # divide by tokens, not samples
loss.backward(); optimizer.step()
```

Defaults. `group_size=16`, `clip_eps_low=0.2`, `clip_eps_high=0.28`, `max_generation_length=20480`, `overlong_cache_length=4096`, `learning_rate=1e-6`, constant learning-rate schedule (per paper). The group size is deliberately larger than GRPO/GSPO/GEPO because dynamic sampling needs candidates to find non-degenerate groups. Optional vLLM acceleration that reuses sampler log-probs.

When to use. Long chain-of-thought reasoning, math/code tasks, scenarios where reward sparsity is the bottleneck.

5.5 VAPO (Value-Augmented Policy Optimization)

Source. Yue et al. [3], ByteDance. Implementation in training/vapo_trainer.py (1036 LOC). The paper reports 60.4 on AIME 2024, we have not independently reproduced this number.

Key innovation. Reintroduce a **value network**, but train it asymmetrically against the policy.

The total loss is a three-term composite:

$$\mathcal{L}_{\text{total}} = \mathcal{L} + c_v \cdot \mathcal{L}_{\text{value}} + w \cdot \mathcal{L}_{\text{positive-LM}}$$

with four supporting mechanisms:

1. **Value-network warmup.** Train the value head alone for 50 steps using Monte Carlo returns before joint training. Mitigates the initialization bias that destabilizes early policy updates.
2. **Decoupled GAE.** The critic uses $\lambda = 1.0$ (unbiased MC); the policy uses a **length-adaptive** λ , computed per-sequence as:

$$\lambda_{\text{policy}} = 1 - \frac{1}{\alpha \cdot |y| + 1}, \quad \alpha = 0.05$$

For a 100-token response this gives $\lambda \approx 0.83$; for a 1000-token response, $\lambda \approx 0.98$. This balances bias and variance differently for short vs. long trajectories.

1. **Positive-example LM loss.** On samples flagged correct by the verifier, add a standard next-token NLL term. This stabilizes the policy when reward signal is sparse: the model is still moving toward the distribution of correct outputs.
2. **Clip-Higher.** Same asymmetric clipping as DAPO ($\text{eps_low}=0.2$, $\text{eps_high}=0.28$).

Implementation citations. ValueHead (2-layer GELU MLP, scalar output) at `vapo_trainer.py:177-223`. LengthAdaptiveGAE at lines 225-349; the exact length-adaptive λ is computed at line 256 as $1 - 1/(\alpha \cdot L + 1)$, evaluated per-sequence (not vectorized across the batch) so that each sequence gets its own λ . `warmup_value_network` at lines 575-680 — trains *only* the value head against MC returns; the policy is frozen for the warmup window. Optional value clipping (PPO-style $V_{\text{clipped}} = V_{\text{old}} + \text{clamp}(V - V_{\text{old}}, \pm \text{clip})$) at lines 729-738.

Defaults. `group_size=16`, `value_warmup_steps=50`, `lambda_critic=1.0`, `lambda_policy_alpha=0.05`, `value_loss_coef=0.5`, `positive_lm_weight=0.1`, `actor_learning_rate=1e-6`, `critic_learning_rate=2e-6`. Separate optimizers and schedulers for actor and critic.

When to use. Highest-quality reasoning training when compute permits. Trades off $\sim 2\times$ memory (value network) and slower convergence (warmup) for SOTA final accuracy.

Pinned to source commit. See Appendix C for reproducibility commands.

5.6 Comparative Summary Summary SL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

Property	GRPO	GSPO	GEPO	DAPO	VAPO
	Token			Token	

Property	GRPO	GSPO	GEPO	DAPO	VAPO
Importance weight		Sequence (length-norm.)	Group expectation		Token + value
Clipping	Symmetric	Symmetric	Symmetric	Asymmetric (Clip-Higher)	Asymmetric
Loss normalization	Sample	Sequence	Group	Token	Token
Advantage baseline	Group	Group	Group	Group	Decoupled GAE
Value network	No	No	No	No	Yes (warmup + decoupled λ)
Defining mechanism	Library baseline	Length normalization	Group expectation	Dynamic sampling + overlong shaping	Length-adaptive GAE + positive LM
Reported AIME score	—	—	—	50/60	60.4
Memory cost	Medium	Medium	Medium	Medium	High
Best for	General	Long outputs, MoE	Async/off-policy	Reasoning	SOTA reasoning

5.7 Exact Forward KL — a quiet technical differentiator

The framework's KL regularizer (`base_trainer.py:600-631` , `compute_kl_divergence`) departs from a shortcut common across the open-source RL stack. There are three patterns in widespread use, and only the first is correct under the policy-gradient signal:

1. **The analytical forward KL** — what we ship. $\sum_v \pi_{\theta}(v) \cdot (\log \pi_{\theta}(v) - \log \pi_{\text{ref}}(v))$ over the full vocabulary at every response position.
2. **Schulman's k_1 approximation** — $0.5 * (\log p - \log q)^2$. The current TRL PPO and RLOO trainers use this (verifiable at [trl/trainer/ppo_trainer.py:517](#) and [trl/trainer/rloo_trainer.py:436](#): `approxkl = 0.5 * (logprobs_diff**2).mean()`). It's a single-sample second-order approximation — fast, low-memory, but biased: the bias is $O(\Delta \log p^3)$ which becomes material once the policy moves more than a few nats from the reference, and it's *strictly non-negative* by construction even when the true KL is small.

3. **Schulman's k_3 estimator** — $r - 1 - \log r$ where $r = \pi/\pi_{\text{ref}}$. Lower-variance than k_1 but still a single-sample Monte Carlo estimator, and several open implementations of k_3 get the direction wrong (computing $\mathcal{KL}(\pi_{\text{ref}} \parallel \pi_{\theta})$ instead of $\mathcal{KL}(\pi_{\theta} \parallel \pi_{\text{ref}})$). PyTorch's `F.kl_div` has this same direction problem if you don't read its signature carefully.

The analytical forward KL we use:

$$\mathcal{KL}(\pi_{\theta} \parallel \pi_{\text{ref}}) = \sum_v \pi_{\theta}(v) \cdot \bigl(\log \pi_{\theta}(v) - \log \pi_{\text{ref}}(v) \bigr)$$

evaluated over the full vocabulary at every response position, then masked to response tokens and normalized by the response length. The cost is one extra `softmax(logits)` per step (already cached when computing log-probs); the benefit is that the KL term is **unbiased and exact**, not an estimator. For small group sizes this materially reduces variance in the regularization signal — exactly when you most need the KL anchor to be reliable. The trade-off is memory: the materialized `current_probs` tensor is $|\text{batch}| \times |\text{seq}| \times |V|$, which can be ~ 1 GB at 32k vocab and 2k sequence length. The framework supports gradient checkpointing of this term for long-context training.

Why this matters for the §10.5 / §11.7 KL-anchor story: the destabilization-to-gibberish failure that §10.5 documents was caused by *no anchor at all* (`beta=0.0`). When the anchor is engaged (`beta > 0`), the difference between "biased k_1 estimator" and "exact analytical forward KL" determines whether the anchor pulls the policy in the right direction or drifts off the SFT manifold over many gradient steps. The framework's choice is the structurally correct one — and it's invisible in the §11.7 result table only because §11.7 doesn't yet probe the regime where it dominates.

6. Training Infrastructure

6.1 BaseTrainer Abstractions

`BaseTrainer` (`training/base_trainer.py`, 853 LOC) is a generic abstract class parameterized by a config type. It owns:

- **Optimizer setup.** AdamW with configurable betas (default 0.9, 0.999) and weight decay; LoRA-aware parameter grouping when `use_pft=True`.
Software parameter grouping when `use_pft=True`.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).
- **Scheduler.** Cosine warmup with configurable warmup ratio (default 0.1); VAPO splits this into separate actor and critic schedulers.

- **Log-prob computation.** Per-token gradient-tracking log-probs with response masking — only response tokens contribute to the loss, prompt tokens are excluded via attention masks.
- **KL divergence.** Forward KL $\pi_{\theta} \cdot \log(\pi_{\theta} / \pi_{\text{ref}})$ against an optional frozen reference model.
- **Gradient clipping.** Norm-based clipping (default `max_grad_norm=1.0`) with statistics tracking — `grad_norm`, `grad_max`, `grad_mean`, `num_zero_grads` — for monitoring vanishing/exploding gradients.

6.2 Model Management

`BaseModelManager` provides unified model loading across all trainers:

- **LoRA / QLoRA** via `peft` (configurable `r`, `lora_alpha`, `lora_dropout`).
- **Quantization** via `bitsandbytes` (4-bit NF4 default, 8-bit option).
- **Mixed precision** (`bfloat16` recommended on H100/A100; `float16` fallback).
- **Gradient checkpointing** for memory-bound training.
- **Reference model** management for KL-regularized objectives.

6.3 Distributed Training

The framework integrates with **Hugging Face Accelerate** for single-node multi-GPU training and **DeepSpeed** (optional `[distributed]` extra) for multi-node training with ZeRO stages 1-3. Ray Train integration is available via the `[hpo]` extra for hyperparameter sweeps.

6.4 vLLM Integration

For trainers that support it (TRL GRPO, DAPO), the rollout phase can be offloaded to vLLM [14]. The trainer's policy weights are synced to a vLLM engine at the start of each rollout round; vLLM generates the batch with paged attention; sampler log-probs are passed back to the trainer to avoid a redundant forward pass.


The speedup is real but variable. The benchmark suite under `benchmarks/performance_benchmarks.py` measures it for a given (model, hardware, batch, sequence-length) combination — we publish the methodology rather than a single multiplier because reported numbers from comparable frameworks span roughly 3× to 20× depending on configuration. Practitioners should measure on their own workload before assuming a specific factor. The dominant variable is batch size: vLLM's edge over HF `generate` widens substantially as `group_size × prompt_batch_size` grows.

Validated configuration. The framework's vLLM integration is validated against Qwen/Qwen2.5-0.5B-Instruct on A100-40GB at `prompt_batch_size=32`, `num_generations=4`, `max_completion_length=512`. Reproduce via [notebooks/vllm_speedup_benchmark.ipynb](#). The result artifact at `benchmark_results/whitepaper_v1/vllm_speedup_qwen25_05b_instruct.json` records the measured (HF, vLLM, ratio) triple

for that configuration plus sweeps over `prompt_batch_size` $\in \{1, 8, 32, 128\}$ so readers can verify the "speedup grows with batch size" claim.

6.5 The Training Data Flow

Every group-based trainer follows the same five-phase loop. The phases are explicit in the codebase and worth understanding because they determine where parallelism is exploited and where serialization is unavoidable:



- **Phase 1** is cheap and CPU-bound — sampling indices from a scenario dataset.
- **Phase 2** dominates wall-clock time in most configurations. This is where vLLM matters most (§6.4); the speedup is workload-dependent.
- **Phase 3** is bottlenecked by reward-function I/O (LLM-judge calls, external verifiers). `asyncio.gather` over the `$G$` responses is the canonical pattern; for hosted LLM judges, batched async concurrency can be the dominant lever on step time.
- **Phase 4** is microseconds — a single tensor op per group.
- **Phase 5** is GPU-bound: forward pass, loss, backward, optimizer step. Memory dominates; speed is secondary.

DAPO adds a **filter step** between 3 and 4 (the `DynamicSamplingBuffer` rejects groups with degenerate accuracy). VAPO adds a **value forward pass** in step 5. GEPO replaces the standard step-4 advantage with the group-expectation coefficient.

When an agent uses the memory subsystem (§3.6), its prompt at step 2 is *constructed* from the current `EnvironmentState.context` plus retrieved long-term and semantic memory. The retrieval happens before generation and serializes into the resulting `ConversationTurn`'s `metadata` so the trainer sees exactly what the policy saw. Memory is not part of the loss — it is, from the trainer's perspective, just part of the conditioning context — but it materially affects what trajectories get sampled in step 2 and what context the reward function in step 3 has access to.

6.6 Offline RL Trainers

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

The framework's online group-based trainers assume a fresh on-policy sampling phase per step. When that's infeasible — e.g., training from logged customer-support

conversations, or warm-starting from a static demonstration set — the framework provides an offline RL track via `OfflineGRPOTrainer` (`training/offline_grpo_trainer.py`). `OfflineGRPOConfig` selects the offline value-learning algorithm via a single field, `offline_algorithm` , with five supported values:

Algorithm	Mechanism
CQL (Conservative Q-Learning)	Adds a conservatism penalty that pushes down Q-values on out-of-distribution actions, preventing overestimation bias.
IQL (Implicit Q-Learning)	Avoids querying out-of-distribution actions entirely by using expectile regression on in-distribution actions. The default; safest under distribution shift.
BCQ (Batch-Constrained Q-Learning)	Restricts the policy to actions close (in a generative-model sense) to the data distribution.
BEAR (Bootstrap Error Accumulation Reduction)	Constrains the policy via MMD distance from the behavior policy.
DT (Decision Transformer)	Sequence-modeling formulation: condition on desired return, predict actions.

The trainer supports a **blended schedule**: an initial offline-only warmup (`warmup_offline_steps=1000` by default) followed by a linear/exponential/constant transition to online GRPO updates. The offline and online value contributions are mixed via `offline_weight` and `online_weight` (both 0.5 by default), and a hybrid baseline (`baseline_type="hybrid"`) averages offline V-values with online group baselines.

State and action embeddings come from a configurable sentence-transformer (`embedding_model="all-MiniLM-L6-v2"` by default), so trajectories of arbitrary text length get fixed-dimensional representations for the value network.

When to use. Bootstrapping from logged conversations; environments where rollouts are expensive (real users, real money); safety-constrained training where you want a behavior-policy regularizer.

6.7 Continual Learning

Skim on first read — opt-in via `[continual]` extra; not on the §11.7 critical path. Read when planning a streaming-data deployment.

Production agents see distribution shift: new product categories, new user demographics, new policies. Catastrophic forgetting — where fine-tuning on new tasks erases prior capabilities — is the canonical failure mode.

Pinned to source commit. See Appendix C for reproducibility commands.
`stateset_agents/training/continual_learning.py` provides seven configurable strategies via a single `strategy` enum:

Strategy	Mechanism	Memory cost
<code>none</code>	Standard fine-tuning baseline	0
<code>replay</code>	Sample from a <code>TrajectoryReplayBuffer</code> of past tasks; mix into current batches	Buffer size
<code>lwf</code>	Learning without Forgetting: KL penalty against a frozen snapshot of the pre-task model	1× model weights
<code>ewc</code>	Elastic Weight Consolidation: per-parameter Fisher-information regularizer pulling toward pre-task weights	1× model weights + Fisher diagonal
<code>replay+lwf</code>	Replay + LwF	Buffer + 1× weights
<code>replay+ewc</code>	Replay + EWC	Buffer + 1× weights + Fisher
<code>replay+lwf+ewc</code>	All three	Maximum

The `TrajectoryReplayBuffer` supports four sampling modes — uniform, recent-biased, reward-weighted, and balanced-by-task — and two storage policies (reservoir for unbounded streams, FIFO for fixed-window). Reward-weighted sampling is particularly useful when high-reward trajectories are rare: they get preferentially preserved across task boundaries.

6.8 The Rust Acceleration Core

A subset of the framework's hottest paths — group-advantage computation, GAE, importance-ratio math, reward normalization — are implemented in Rust and exposed to Python via **PyO3**. The crate `stateset-rl-core` lives at `rust_core/` and compiles as both `cdylib` (for Python import) and `rlib` (for downstream Rust integration).

Functions exposed to Python:

Rust function	Purpose
<code>compute_group_advantages</code>	Within-group reward normalization with Welford's online algorithm
<code>compute_gae</code>	Single-trajectory Generalized Advantage Estimation with configurable γ , λ
<code>batch_compute_gae</code>	Parallel GAE across a batch of trajectories, parallelized via Rayon
<code>compute_gspo_importance_ratios</code>	Length-normalized sequence ratios for GSPO
<code>compute_ppo_clipped_surrogate</code>	The standard PPO clipped objective
<code>normalize_with_running_stats</code>	

Rust function	Purpose
	Streaming Welford normalization for cross-batch advantage statistics

The performance rationale: Python's overhead dominates when the inner loop isn't NumPy-vectorizable. GAE's backward recurrence is the canonical example — each step depends on the next, defeating BLAS — and that's where the Rust core's advantage compounds with batch size.

Read this before the speedup table below. The numbers in the next table are *isolated-kernel* microbenchmarks. They are not a framework-level speedup claim. In an actual §11.7 training step on Colab A100, Phase 2 (generation) consumes roughly 70–85% of wall-clock (\$6.5), the GAE/advantage Phase 4 is ~1–3%, and Phase 5 forward+backward is ~10–20%. **A 30× speedup on Phase 4 translates to <1% improvement in end-to-end step time on this configuration.** The Rust core is the right primitive for Phase-4-dominated workloads (very large group sizes, long trajectories, value-net actor-critic training) and for keeping the recurrence-heavy paths *off* the bottleneck heat map. It is *not* a blanket framework speedup. Be honest about this with your profiler: a per-step `time.perf_counter()` around each kernel will tell you which side is the bottleneck on your hardware.

Measured speedup (`benchmark_results/whitepaper_v1/rust_vs_python_microbenchmark.json` , single-process CPU microbenchmark, float64):

Function	Batch × T	Python-loop	Rust (Rayon)	Kernel speedup	End-to-end §11.7 impact
<code>batch_compute_gae</code>	8 × 256	1.49 ms	0.06 ms	26×	<1%
<code>batch_compute_gae</code>	32 × 512	7.05 ms	0.10 ms	72×	<1%
<code>batch_compute_gae</code>	128 × 1024	52.8 ms	1.51 ms	35×	<1%

For *vectorizable* kernels (e.g. `compute_group_advantages` on a 2-D ndarray), NumPy with BLAS is competitive with the Rust path — at typical batch sizes the function-call boundary and ndarray conversion overhead can leave Rust slightly behind plain NumPy. The Rust core is included for the recurrence-heavy paths (GAE, importance-ratio backward sums) — *not* as a blanket NumPy replacement.

The Rust core is **optional**. The framework includes pure-Python fallbacks for every accelerated function. CI builds the Rust core via `maturin` and tests both paths; users who can't compile Rust still get a working framework, just slower on GAE.

Configuration in `Cargo.toml`: release builds use `opt-level=3`, `LTO=fat`, `codegen-units=1`, `panic=abort` — standard Rust release settings.

Placed to some credit. See `build.rs` for some details. `CC0` or `MIT` license. Rust 1.1 (transitioning to Apache 2.0 on 2029-09-03).

6.9 Sim-to-Real Transfer

Skim on first read — relevant if you're bridging templated/scripted user models to real user populations. The trainer doesn't require either module.

Most production agents are trained against *simulated* users (templated personas, scripted user models) and then deployed against real ones. The simulator-to-reality gap is real: real users are more variable, less cooperative, less articulate.

The framework provides two complementary modules:

- `domain_randomization.py` — randomizes the simulator. A `UserPersona` has seven core traits (patience, expertise, verbosity, formality, emotional stability, cooperativeness, detail orientation) plus response-level noise (typos, truncation, latency simulation). Five pre-defined templates cover canonical edge cases: `patient_expert`, `frustrated_novice`, `busy_professional`, `curious_learner`, `skeptical_critic`. Curriculum learning ramps difficulty from 0.3 → 0.9 over a configurable horizon (default 5,000 steps).
- `sim_to_real.py` — bridges sim and real. A `UserBehaviorModel` learns to predict real-user responses (embedding, length, emotion, continuation probability) from logged conversations. A `DomainAdaptationModule` then aligns sim and real distributions using one of three methods: **DANN** (adversarial domain confusion), **MMD** (maximum mean discrepancy minimization), or **CORAL** (covariance alignment). A progressive-transfer schedule decays the simulation weight from 1.0 to 0.1 over training, with early stopping based on monitored sim-real gap.

These modules are designed to compose with any trainer: the simulator is just a `ConversationEnvironment` subclass, so DAPO or VAPO can run on top of a domain-randomized user model with no algorithmic change.

7. Operational Layer

7.1 FastAPI Service

`stateset_agents/api/main.py` exposes the framework as an HTTP service. Key endpoints:

Path	Purpose
<code>POST /v1/messages</code>	Multi-turn conversation (Anthropic-compatible)
<code>POST /v1/chat/completions</code>	OpenAI-compatible chat
<code>POST /agents</code> · <code>GET /agents/{id}</code>	Agent lifecycle
<code>POST /agents/{id}/messages</code>	Single-turn inference
<code>POST /training</code> · <code>GET /training/{id}</code>	Training job submission and status

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

Path	Purpose
GET /metrics	Prometheus metrics
GET /ready · /live	Kubernetes health probes
GET /circuits	Circuit breaker status

A `LazyApp` ASGI proxy defers application initialization to the first request, keeping cold-start latency low. The lifespan manager initializes a configurable cache backend (in-memory or Redis), an optional persistence backend (PostgreSQL or SQLite), an `AgentService`, and a `TrainingService`. Middleware stack: GZip, CORS, rate limiting, security headers.

Example: send a message to an agent.

```
curl -X POST https://api.example.com/v1/messages \
-H "Authorization: Bearer $TOKEN" \
-H "Content-Type: application/json" \
-d '{
  "model": "agent-customer-support-v3",
  "messages": [
    {"role": "user", "content": "My order #12345 hasn't arrived"}
  ],
  "max_tokens": 512,
  "temperature": 0.7
}'
```

Example: kick off a training job.

```
curl -X POST https://api.example.com/training \
-H "Authorization: Bearer $ADMIN_TOKEN" \
-H "Content-Type: application/json" \
-d '{
  "algorithm": "gsdp",
  "base_model": "Qwen/Qwen3.5-0.8B",
  "environment": "customer_support_v2",
  "config": {"group_size": 4, "learning_rate": 5e-6, "num_epochs": 3}
}'
# → {"training_id": "tr_alb2c3", "status": "queued"}
```

Example: poll a training job.

```
curl https://api.example.com/training/tr_alb2c3 \
-H "Authorization: Bearer $ADMIN_TOKEN"
# → {"status": "running", "step": 142, "loss": 0.0341, "reward_mean": 0.62, ...}
```

The OpenAI-compatible `/v1/chat/completions` endpoint accepts the same request shape as the OpenAI API, which means clients written against OpenAI (Python SDK, LangChain, LiteLLM, OpenWebUI) work against a StateSet Agents service with only a base-URL change.

7.2 Deployment Artifacts

Docker. A multi-stage Dockerfile compiles the Rust acceleration layer from source, then ships a Debian-slim runtime under a non-root user. The bundled `docker-compose.yml` provides a full local stack: the API, the StateSet commerce backend (`stateset-api` , pinned tag), PostgreSQL 16, Redis 7.

Kubernetes. A Helm chart (v0.1.0) ships with 10+ values overlays for different GPU profiles (A100, H100, B200, Kimi-K2.5 fine-tuned, GLM 5.1 FP8). GKE-specific overlays cover both Autopilot and Standard cluster types, with separate staging/production examples. Raw manifests under `deployment/kubernetes/` cover training jobs (Qwen 3.5 27B, Kimi-K2.5, GLM 5.1) and vLLM serving deployments.

7.3 Observability

A production training and serving stack must be observable. The framework exposes two metric surfaces.

HTTP-layer Prometheus metrics (`stateset_agents/api/middleware.py`):

Metric	Type	Labels
<code>stateset_http_requests_total</code>	Counter	<code>method</code> , <code>endpoint</code> , <code>status_code</code>
<code>stateset_http_request_duration_seconds</code>	Histogram	<code>method</code> , <code>endpoint</code>
<code>stateset_http_requests_in_progress</code>	Gauge	<code>method</code> , <code>endpoint</code>

These are the standard latency/throughput/error triplet — the basis of any RED-method ([Rate](#), [Errors](#), [Duration](#)) dashboard.

Domain-specific metrics (`stateset_agents/api/grpo/metrics.py`): a `GRPOMetrics` class tracks:

- Request counts and status counts (per endpoint)
- Latencies (a rolling `deque` , with `get_latency_percentiles()` computing avg/p50/p95/p99)
- Error counts and rate-limit hit counts
- Training-job lifecycle: `training_jobs_started` , `training_jobs_completed` , `training_jobs_failed` , `total_trajectories` , `total_computation`
- Conversation lifecycle: `conversations_started` , `conversations_ended` , `messages_processed`
- WebSocket: `websocket_connections` , `websocket_messages`

These domain metrics are exposed at:

Pinned to source commit. See Appendix C for reproducibility commands.

- `GET /metrics` — Prometheus text format (scrape from Prometheus / Grafana Agent / OpenTelemetry Collector).
- `GET /metrics/json` — admin-authenticated JSON dump (useful for ad-hoc debugging).

- `GET /metrics/security` — security-event metrics (rate-limit hits, auth failures).
- `GET /metrics/cache` — cache hit/miss/size.
- `GET /health` — composite health check, cached for 5 seconds to avoid hammering downstreams.

Prometheus export is gated by the `API_ENABLE_PROMETHEUS` environment variable so the dependency is opt-in.

7.4 The Dashboard

A React 19 + Vite + TypeScript dashboard ships in `dashboard/` for interactive monitoring and experiment management. It is a thin client over the FastAPI service — no direct database access — and provides six views (screenshots in [dashboard/README.md](#); run locally with `make dashboard` to see the UI directly):

1. **Dashboard** — overview of recent experiments
2. **Create Experiment** — submit a new training run with a configuration form
3. **Live Monitor** — stream loss/reward curves from an active run (polls `/experiments/{id}/metrics` every 5 s; Recharts for plotting)
4. **Compare Experiments** — side-by-side metric comparison across runs
5. **Playground** — interactive chat with a trained agent for qualitative testing
6. **Leaderboard** — ranked view of recent runs by objective metric

Keyboard shortcuts (`Cmd+N`, `D`, `P`, `L`, `C`) navigate views; the dashboard is bundled with the API container or can be served independently as a static site.

7.5 Benchmark Methodology

The repository ships a benchmark suite under `benchmarks/` whose output is the basis for any performance claim we make. Rather than publish numbers that will go stale with the next driver, model release, or kernel update, this whitepaper publishes the *methodology* and points to the suite:

- `benchmarks/performance_benchmarks.py` (1128 LOC) — end-to-end training and inference latency, throughput, memory footprint.
- `benchmarks/algorithm_comparison.py` (569 LOC) — head-to-head GRPO vs. GSPO vs. DAPO on a fixed environment and reward.
- `benchmarks/framework_comparison.py` (526 LOC) — StateSet Agents vs. baseline TRL on matched configurations.
- `benchmarks/real_performance_benchmarks.py` (443 LOC) — vLLM vs. plain HF generation, measured at multiple batch sizes.

The shared `BenchmarkResult` dataclass captures: `name`, `iterations`, `avg / min / max / p50 / p95 / p99` latency (ms), `std_dev`, `throughput` (ops/sec), `memory_mb`. Results land in `benchmark_results/` as JSON and (optionally) HTML reports via `--report html`. We recommend practitioners run the relevant slice on their target hardware before

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

committing to a configuration — the headline multipliers vary too much across deployments to publish a single canonical number.

7.6 CI/CD

GitHub Actions (`.github/workflows/`) runs:

- **Test matrix** — Python 3.10/3.11/3.12/3.13 on Ubuntu; 3.10/3.13 on Windows.
- **Lint and type checks** — `ruff`, `black`, `isort`, `mypy` (strict mode gated per-module).
- **Tests** — `pytest` with Codecov upload. Coverage reporting reflects only paths exercised by the in-process unit and integration tests (~49% of total LOC). End-to-end serving, Helm-rendering, and GPU-only training paths are tested but don't contribute to per-line coverage; the headline number understates the tested surface. The CI gate requires the core-abstractions modules to stay above a stricter threshold via Codecov component scoping.
- **Security scans** — `bandit` and `safety` SBOM generation.
- **Helm validation** — `helm lint` and template rendering across all values overlays.
- **Docs** — Sphinx build with RTD theme.

8. Supported Models

StateSet Agents ships first-class starters (with dedicated CLI commands, configuration modules, and three preconfigured profiles each) for:

Model	HF ID	CLI
Qwen 3.5 0.8B	<code>Qwen/Qwen3.5-0.8B</code>	<code>stateset-agents qwen3-5-0-8b</code>
Gemma 4 31B IT	<code>google/gemma-4-31b-it</code>	<code>stateset-agents gemma-4-31b</code>
Kimi-K2.6	<code>moonshotai/Kimi-K2.6</code>	<code>stateset-agents kimi-k2-6</code>
GLM 5.1 (754B MoE)	<code>zai-org/GLM-5.1</code>	<code>module + example (QLoRA-only)</code>

Reference models with examples and Kubernetes manifests: Qwen 3.5 27B, Qwen 3, Qwen 2.5 3B Instruct, Kimi-K2.5, Gemma 3 / 2 27B IT, Llama 3, Llama 2 7B, Mistral 7B, GPT-2. Generic support: any Hugging Face causal LM compatible with `AutoModelForCausalLM` and TRL GRPO.

8.1 Component Maturity

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

Not every module in the framework is at the same level of production-readiness. The table below disambiguates **stable** (used in production deployments, API-stable across

point releases), **beta** (functionally complete, API may change, used in non-critical production), and **experimental** (works in tests, not yet recommended for production).

Component	Maturity	Notes
Agent , MultiTurnAgent , ToolAgent	Stable	Core abstractions; covered by integration tests
ConversationEnvironment , RewardFunction , CompositeReward	Stable	
ModelBackend Protocol + StubBackend / HuggingFace backend	Stable	
vLLM backend	Beta	Sync semantics for policy-weight reloading are still hardening
TRL GRPO trainer	Stable	Delegates to TRL; matches upstream stability
GSPO trainer	Beta	Heavily tested; first-party three-seed positive transfer result documented in §11.7
DAPO trainer	Beta	Same
GEPO trainer	Beta	Same
VAPO trainer	Experimental	Largest surface area; warmup and decoupled-GAE paths still being tuned
Offline GRPO (CQL/IQL/BCQ/BEAR/DT)	Experimental	Each algorithm tested individually; blending schedule is new
Continual learning (replay/LwF/EWC)	Experimental	Lacks the dialogue-specific evaluation harness called out in §10.5
Sim-to-real, domain randomization	Beta	Used internally; public benchmark suite pending
Memory subsystem	Beta	In-memory and Redis backends stable; SQLite less exercised
Neural reward trainer	Beta	
RLAIF trainer	Beta	
Auto-research loop	Experimental	Useful but rapidly evolving API
Rust acceleration core	Beta	Optional; pure-Python fallbacks always present
FastAPI service, observability	Stable	

Component	Maturity	Notes
Helm chart, K8s manifests	Beta	Stable for the GPU profiles we ship; custom profiles need adaptation
Dashboard	Beta	

Treat this matrix as part of the public contract: changes to **stable** components follow semantic versioning; **beta** components may change their config schemas in minor releases (with migration notes); **experimental** components may change their public API in patch releases.

Note on the apparent inversion — the operational layer (FastAPI service, observability) is Stable while the trainers are Beta / Experimental . This reflects the framework's deployment history rather than any value judgment: the serving layer hardened first against StateSet's own commercial customer-support deployments, where API stability is a hard requirement. The trainer surface is currently in active hardening — the v0.13 release line is closing first-party benchmark gates (see §11.7) ahead of promoting any of GSPO/DAPO/GEPO to Stable in v0.14. VAPO remains Experimental because the value-net warmup + decoupled-GAE tuning surface is still being validated against §5.5's reproduced numbers.

9. Testing Philosophy

A core engineering principle of StateSet Agents is that **conversational RL infrastructure should be testable without a GPU**. The framework achieves this via three patterns:

1. **Stub backend.** `StubBackend` provides deterministic responses without any model weights. Tests parameterize agents with `AgentConfig(use_stub_model=True)` and never touch Hugging Face.
2. **Property-based stub detection.** Agents expose `_is_stub_backend` as a property that checks `isinstance(self.model, StubModel)`. Inference services expose `is_stub` analogously. This avoids brittle string-matching on model names.
3. **Canonical exception tuples.** All retry/fallback logic catches from a small set of canonical exception tuples in `stateset agents/exceptions.py` (`IMPORT_EXCEPTIONS`, `GPU_EXCEPTIONS`, `MODEL_IO_EXCEPTIONS`, `INFERENCE_EXCEPTIONS`, `ATTRIBUTE_VALUE_EXCEPTIONS`, `NETWORK_EXCEPTIONS`, `SERIALIZATION_EXCEPTIONS`, `MODEL_DEVICE_EXCEPTIONS`). This makes the failure surface explicit and testable.

The test suite contains 2,438 tests organized into `tests/unit/`, `tests/integration/`, `tests/api/`, `tests/e2e/`, and `tests/performance/`, with `pytest-benchmark` regression tests gating reward throughput and manifest build time. Helm chart rendering is smoke-tested against every values overlay (skipping gracefully if `helm` is not installed).

10. Discussion

10.1 Why Group-Based Methods Win for Dialogue

The fundamental claim of this framework — and of the research it operationalizes — is that **group-based, baseline-normalized policy gradients are the right primitive for conversational RL**. Three reasons:

1. **Variance reduction without a critic.** A separate value network is one of the largest sources of engineering complexity in PPO (sync, scheduling, target networks). Group-relative advantages give you variance reduction "for free" from the group itself, and only VAPO reintroduces a value network — for the marginal accuracy bump in highly demanding reasoning tasks.
2. **Natural fit for multi-turn rollouts.** Sampling G trajectories per prompt is computationally well-suited to batched generation engines (vLLM, TGI). The trajectory is the unit of work, which aligns with how production inference systems actually operate.
3. **Stable on long sequences.** As GSPO showed, sequence-level importance weighting eliminates the multiplicative variance explosion of token-level PPO. This is the difference between a training run that completes and one that diverges at 50% through.

10.2 When to Choose Which Trainer

A pragmatic flowchart:

- **First training run on a new task?** Start with **TRL GRPO**. It's the simplest, best-supported, and benefits from the vLLM acceleration path.
- **Outputs are long (CoT, code, structured generation) or you're using an MoE model?** Switch to **GSPO**.
- **Training is distributed or off-policy?** Use **GEPO**.
- **Reasoning task with sparse reward (math, code competitions)?** Use **DAPO**.
- **You have compute headroom and need SOTA reasoning performance?** Use **VAPO**.

Copyright © 2024 StateSet AI, Inc. All rights reserved. Apache 2.0 License.

Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

10.3 What's Hard About This

Multi-turn conversational RL has structural difficulties that single-turn RLHF can ignore:

1. **Credit assignment across turns.** Reward arrives at the end of a conversation (problem resolved? user satisfied?), but the agent's *individual decisions* are distributed across turns. Group-relative advantages handle inter-trajectory credit but not intra-trajectory credit; that's still an open problem.
2. **The user-model gap.** Training against a simulated user is fundamentally different from training against a real one. The framework mitigates this with domain randomization and sim-to-real transfer (§6.9), but the residual gap is the largest source of deploy-time surprise.
3. **Reward gaming.** Composite rewards encode an *implicit* utility function. Policies are excellent at finding the cheapest way to maximize a literal reward — often producing outputs that score high but feel wrong. The framework's bias toward composite rewards (rather than a single learned scalar) is a partial mitigation; the deeper fix is iteration on the reward function itself. *Inverse case (rule-based-rubric blindness)*: the bundled customer-support benchmark's keyword-presence rubric scores a coherently-trained policy *lower* than its untuned baseline (0.34 vs 0.54) because the trained model learned to pivot to clarifying questions instead of dropping rubric keywords verbatim. The trained model is qualitatively better — see the live-demo evidence in `benchmark_results/whitepaper_v1/customer_support_qwen3_5_0_8b_gspo_klanchor.json` — but the rubric can't see it. Rule-based rewards are bias-stable but blindness-stable in opposite directions; a paraphrase-tolerant LLM-judge is the natural complement.
4. **Online + multi-turn + tool use.** Each of these is a hard problem in isolation; their interaction is harder. Tool calls add hidden state and stochastic latency; multi-turn adds long-horizon credit assignment; on-policy RL adds non-stationarity.

10.4 When StateSet Agents Is the Wrong Choice

Honest scoping matters more than feature lists. StateSet Agents is built for one shape of problem; for other shapes, other tools are better:

- **Single-turn preference data, no environment.** If you have a static dataset of (prompt, chosen, rejected) tuples and no rollout-time interaction, **DPO** [11] in TRL is simpler, has no rollout phase, and converges faster. The group-based machinery here is overhead you won't recover.
- **Pure supervised fine-tuning.** If you don't need RL at all — just SFT on demonstrations — use Hugging Face's `Trainer` or `axolotl` directly. This framework's stub backend and reward machinery add nothing to that workflow.
- **Sub-billion-parameter models on edge devices.** The framework assumes server-class GPUs and Hugging Face checkpoints. For on-device tiny models with custom runtimes (`ggml` , `MLX` , `ExecutorTorch`), the abstractions here don't compose cleanly.
- **Image, audio, or video modalities.** The agent and environment abstractions are text-first; multimodal trajectories would require substantial extension.

- **Real-time RL with strict latency budgets.** The async-everywhere design is throughput-optimized, not tail-latency-optimized. If you need bounded p99 inference under 50 ms in production, vLLM directly + a thin wrapper is the right shape.
- **Hard-real-time safety-critical systems.** This framework's exception philosophy is graceful degradation; for safety-critical RL you want fail-stop and formal verification, which is not on the roadmap.

These are not failings — they are scope choices. The framework is built for the case where multi-turn interaction matters, the model lives in PyTorch on a GPU, and you can iterate the reward function until it's right. That case is large enough to be worth a dedicated tool.

10.5 Safety and Threat Considerations

A framework for training online conversational agents has a real safety surface. We've made design choices to mitigate the most common risks, but practitioners should understand the residuals:

- **Reward hacking.** All on-policy RL is susceptible to policies that find unintended ways to maximize reward. The framework's bias toward **composite rewards** (multiple components combined with `weighted_sum`) makes this harder than a single learned scalar — gaming one component often hurts another. But it does not eliminate the risk. Operational guidance: always include a `SafetyReward` component, periodically inspect raw generations from the latest checkpoint, and treat rapid reward climbs (especially without corresponding eval-set improvement) as a flag, not a victory.
- **Unsafe exploration.** Group-based methods sample G trajectories per prompt; in production-facing deployment, this means G user-visible outputs per query during training. Do not train against real users without (a) a `SafetyReward` filter on the response distribution, (b) a hard length / content cap before the user sees anything, and (c) a circuit breaker that aborts a training run if any safety metric crosses threshold. The framework provides the hooks; the policy decisions are yours.
- **Data privacy and logging.** Conversation trajectories are PII by default. The `ConversationMemory` Redis/SQLite backends and the `total_trajectories` Prometheus counter both touch user content. Operators should: (a) configure log retention to match their privacy posture, (b) ensure trajectory data is encrypted at rest in any persistence backend, (c) scrub PII before any cross-region replication, and (d) maintain a clear data-deletion path for user erasure requests.
- **Reference-model drift.** The reference model anchors the KL regularizer. If it's the wrong reference (e.g., an outdated SFT checkpoint), the regularizer pulls toward outdated behavior. Re-anchor the reference after major fine-tuning rounds. **Absence is worse than drift.** Running GSPO with `use_reference_model=False` and `beta=0.0` on a small corpus is the canonical "policy goes off the rails" setup. A documented first-party case is in `benchmark_results/whitepaper_v1/customer_support_qwen3_5_0_8b_gspo.json` — 3 epochs over 16 scenarios under the

bundled defaults destabilized the model to the point of emitting token soup, while a single-variable fix (`use_reference_model=True`, `beta=0.05`) restored coherent customer-service English (`customer_support_qwen3_5_0_8b_gspo_klanchor.json`). `train_with_gspo` now emits a runtime warning when the unsafe combination is detected on small corpora.

- **Tool-use blast radius.** `ToolAgent` invokes tools with real-world side effects. During training, restrict the registered tool set to read-only or sandboxed variants; only enable write/destructive tools after the policy has demonstrated stable tool-call patterns in evaluation. The framework does not enforce this — it's a deployment-time decision.
- **Judge-model power asymmetry.** RLAIIF (§4.5) bootstraps from a stronger judge. If the judge model itself has biases or failure modes, the policy will inherit them. KL regularization against a frozen reference (`beta > 0`) is the primary mitigation; periodic human spot-checks of high-reward / low-reward judge calls are the secondary one.

None of these issues are unique to this framework, but the multi-turn online setting amplifies them. We treat safety as an operational concern that needs both framework-level affordances (which we provide) and deployment-level discipline (which we cannot enforce).

10.6 Open Questions and Future Work

- **Hierarchical credit assignment** across turns: current trainers assign group-relative credit per trajectory but do not attribute reward to specific turns. Per-turn baselines and turn-level group normalization are natural extensions.
- **Continual learning under distribution shift:** the framework includes EWC, LwF, and replay primitives (§6.7) but lacks a unified evaluation harness for catastrophic forgetting in dialogue. A standard benchmark would unlock systematic comparison.
- **Reward-model uncertainty:** composite rewards aggregate point estimates; propagating uncertainty (e.g., Bayesian reward heads or ensembles) could improve robustness to reward-model errors and provide a principled exploration signal.
- **Tool use as a first-class RL signal:** `ToolAgent` supports tool invocation, but tool-execution rewards are currently composed externally rather than integrated into a unified RL objective with credit assignment to the tool-call decision itself.
- **Asynchronous, multi-actor training:** GEPO is the framework's nod in this direction, but a complete actor-learner separation (Ape-X / Impala style) for LLM agents is an unsolved engineering problem at the framework level.

11. Practitioner's Guide

11.1 Quick Start

The minimal end-to-end training loop, using GSPO on a small Qwen model with a custom domain reward:

```
import asyncio
from stateset_agents.core import MultiTurnAgent, AgentConfig
from stateset_agents.core import ConversationEnvironment
from stateset_agents.core.basic_rewards import HelpfulnessReward
from stateset_agents.core.reward_base import CompositeReward
from stateset_agents.training import GSPOTrainer, GSPOConfig

async def main():
    agent = MultiTurnAgent(AgentConfig(
        model_name="Qwen/Qwen3.5-0.8B",
        torch_dtype="bfloat16",
        attn_implementation="flash_attention_2",
        use_peft=True,
        peft_config={"r": 16, "lora_alpha": 32, "lora_dropout": 0.05},
    ))
    await agent.initialize()

    reward_fn = CompositeReward(
        components=[HelpfulnessReward(weight=1.0)],
        combination="weighted_sum",
    )

    env = ConversationEnvironment(
        scenarios=load_scenarios("data/customer_support.jsonl"),
        max_turns=8,
        reward_fn=reward_fn,
    )

    trainer = GSPOTrainer(
        config=GSPOConfig(
            group_size=4,
            clip_range_left=3e-4,    # See §5.2: GSPO ratios are length-normalized,
            clip_range_right=4e-4,   # so the effective clip must be much tighter than PP0's 0.2.
            learning_rate=5e-6,
            num_train_epochs=3,
        ),
        agent=agent,
        environment=env,
    )
    await trainer.train()

asyncio.run(main())
```

A test-only equivalent — runs without any model weights, GPU, or network access — is identical except for one config line:

```
agent = MultiTurnAgent(AgentConfig(use_stub_model=True))
```


This is the seam that lets the same training script run in CI as runs on an H100 cluster.

11.2 Custom Reward Functions

Rewards are async and inherit from `RewardFunction`:

```
from stateset_agents.core.reward_base import RewardFunction, RewardResult
from stateset_agents.core.trajectory import ConversationTurn

class JSONValidityReward(RewardFunction):
    name = "json_validity"
    weight = 0.5

    async def compute_reward(
        self,
        turns: list[ConversationTurn],
        context: dict | None = None,
    ) -> RewardResult:
        last = turns[-1].content
        try:
            import json; json.loads(last)
            return RewardResult(score=1.0, breakdown={"json_valid": 1.0})
        except Exception as e:
            return RewardResult(
                score=0.0,
                breakdown={"json_valid": 0.0},
                explanation=f"Invalid JSON: {e}",
            )
```

Wrap multiple rewards in a `CompositeReward` to get a single trainer-compatible signal.

11.3 Choosing a Trainer: A Decision Tree

```
Start
|
├─ First training run on this task? → TRL GRPO
|
├─ Long outputs (CoT/code) or MoE model? → GSP0
|
├─ Off-policy / replay / async training? → GEPO
|
├─ Sparse-reward reasoning task? → DAP0
|
└─ Need maximum reasoning accuracy, have GPU headroom? → VAP0
```

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

What this decision tree is and isn't. The branches summarize the *literature* on each trainer's design intent (GSPO paper's length-normalization argument, DAPO paper's Clip-Higher + dynamic-sampling story, etc.) plus our own implementation experience. First-party head-to-head numbers across these trainers on a shared task are pending — the [notebooks/whitepaper_v1_comparative_trainers.ipynb](#) protocol is the harness, but the result row isn't in §5.6 or §11.7 yet. Use this tree as starting advice; replace any branch with measured guidance once your own §5.6 comparison run lands.

11.4 Operational Recipes

Goal	Setting
Speed up rollouts (typically several × on large batches)	Install <code>[vllm]</code> extra; set <code>use_vllm=True</code> in trainer config
Train a 30 B+ model on a single 80 GB GPU	LoRA + bitsandbytes 4-bit quantization + gradient checkpointing
Multi-node training	<code>[distributed]</code> extra → DeepSpeed ZeRO-3
Hyperparameter search	<code>[hpo]</code> extra → Optuna with the built-in <code>objective()</code> template
Deploy as a service	<code>[api]</code> extra → <code>uvicorn stateset_agents.api.main:app</code> ; Helm chart for K8s
Train without GPU (CI, prototyping)	<code>AgentConfig(use_stub_model=True)</code>

11.5 Failure Modes and Diagnostics

Group-based RL has characteristic failure modes that look different from supervised fine-tuning. This subsection collects the most common ones we've seen, their signatures, and the first thing to check.

Symptom	Likely cause	First thing to check
Loss → 0, reward flat	The clip range is too tight; gradients are vanishing because every sample is clipped. Pinned to source commit. See Appendix C for reproducibility commands.	For GSPO, widen <code>clip_range_left</code> / <code>clip_range_right</code> . For DAPO/VAPO, raise <code>clip_eps_high</code> .
Loss → 0, reward → 0	Dynamic sampling is filtering out the whole batch. Licensed BUI SL 1.1; (transitioning to Apache 2.0 on 2029-09-03)	DAPO: check <code>min_accuracy_threshold</code> / <code>max_accuracy_threshold</code> ; verify the reward function actually produces a spread of values.

Symptom	Likely cause	First thing to check
Reward climbs then collapses	Mode collapse — policy has found a degenerate response that scores high under the reward.	Inspect generations directly; add a length or diversity penalty to the composite reward.
Gradient norms explode	Importance ratios going unstable.	Lower learning rate. For long sequences, prefer GSPO over GRPO.
Gradient norms vanish	The policy is already saturating the reward, or KL is dominating the loss.	Lower <code>beta</code> ; check whether the policy is near-deterministic (<code>top_p</code> too low?).
OOM on H100	Reference-model + base-model + KV cache + activations exceed memory.	Enable <code>gradient_checkpointing=True</code> ; switch to QLoRA with <code>use_4bit=True</code> ; reduce <code>max_completion_length</code> .
vLLM speedup not realized	Generation is <i>not</i> actually the bottleneck — the reward function is.	Profile each phase (§6.5). Move LLM-judge rewards to a faster judge model or pre-compute via the neural-reward path (§4.4).
Training works on stub, fails on HF	A test-only assumption leaked into production code.	Verify <code>_is_stub_backend</code> checks are not gating production behavior; check that tokenization handles real chat templates.
Loss spikes at episode boundaries	Multi-turn context not being reset properly between episodes.	Confirm <code>await agent.reset()</code> is being awaited (it's async — see Project Memory).
Reward function raises but training continues	Default <code>CompositeReward</code> graceful-degradation behavior; one component is failing silently.	Check W&B for per-component reward breakdowns; failures log warnings but don't abort.
Trained model emits token soup; rubric score still nonzero	No KL anchor + small corpus + rule-based reward → policy drifts off the coherent text manifold while still hitting rubric keywords.	Set <code>use_reference_model=True</code> , <code>beta=0.05</code> ; <code>train_with_gspo</code> emits a runtime warning when this combination is detected on a small corpus. See §10.5 and <code>benchmark_results/whitepaper_v1/customer_support_gwen3.5-90b_gspo.json</code> .

🚩 The token-soup failure deserves its own callout. If the rubric score is nonzero but the live-demo output is gibberish, that's a *coherence collapse*, not a training success. The trained policy has found a way to hit rubric keywords without producing intelligible English. This is the canonical "RL without KL anchor on small corpus" failure mode. Never accept a rubric improvement without spot-checking the actual generations. §11.7's protocol does this via the LLM-judge eval; the keyword rubric alone cannot distinguish "good response that mentions refund" from "random tokens that contain the word 'refund'."

Key W&B / metric panels to watch:

- `reward/mean` , `reward/std` — primary signal. `std` collapsing toward 0 within a group means the policy has converged to a single response per prompt — usually bad.
- `advantage/mean` , `advantage/std` — should hover around (0, 1) after normalization. Otherwise something is wrong with reward scaling.
- `policy/clip_fraction` — for token-level methods, healthy values are 0.05–0.20. Above 0.5 means clipping is dominant and effective step size is small.
- `policy/entropy` — collapsing entropy means the policy is becoming deterministic. Some collapse is expected; falling off a cliff is not.
- `kl/forward` — if $\beta > 0$, this should grow slowly. Sudden jumps mean the policy is drifting hard from the reference.
- `gradient_stats/grad_norm` , `grad_max` , `num_zero_grads` — `BaseTrainer` tracks these; spikes are precursors to instability.

11.6 Autonomous Research Loops

Skim on first read — opt-in via `[auto-research]` extra. Useful for overnight HPO sweeps, not required for any §11.7 result.

For practitioners who want to *automate* the hyperparameter search itself, the framework ships `training/auto_research/` — a self-driving experiment loop. The core `AutoResearchLoop` class orchestrates four phases on a repeating schedule:

1. **Propose** a new configuration via a pluggable `ExperimentProposer` .
2. **Train** the agent with a wall-clock time budget enforced by `asyncio.wait_for` .
3. **Evaluate** on a held-out scenario set.
4. **Decide** whether to keep the proposal as the new incumbent or revert.

Four proposer strategies are built in:

Proposer	Pinned to source commit. See Appendix C for reproducibility commands. Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).	
	Strategy	
<code>RandomProposer</code>	Uniform sampling over a declared search space	
<code>PerturbationProposer</code>	Gaussian perturbation around the current incumbent	

Proposer	Strategy
BayesianProposer	Surrogate-model-driven exploration (Gaussian-process-style acquisition)
LLMProposer	Prompts an LLM (Claude by default) with the run history and asks for the next config

ExperimentTracker logs all runs with their objective metric and direction (min/max).

CheckpointManager persists artifacts to the filesystem (no git dependency).

EarlyAbortManager monitors convergence per-run and aborts trials that aren't improving — saving compute on dead-end configurations.

The whole module is loaded via the optional [auto-research] extra. A typical use is overnight HPO sweeps over learning rate, KL coefficient, group size, and LoRA rank for a fixed environment + reward function. The LLM proposer is particularly valuable for irregular, non-numeric search spaces (e.g., "should we enable positive-LM loss?").

11.7 First-Party Reproduction (Customer Support, GSPO)

This is the **canonical first-party benchmark** for the v1.0 whitepaper. Result file:

[benchmark_results/whitepaper_v1/](#)

[customer_support_3seed_judge_qwen25_05b_instruct.json](#). Reference notebook: [notebooks/customer_support_3seed_judge.ipynb](#). Hardware: Colab A100-40GB. Wall clock: ~11 min for three-seed training + dual eval.

	Baseline	Trained (mean)	Δ	σ across seeds	3-seed agreement
Composite rubric (keyword-based)	0.506	0.578	+0.072	0.067	✓
LLM-judge (paraphrase-tolerant)	0.750	0.829	+0.079	0.058	✓

Methodology. Three seeds: 42, 1337, 2026. Trainee: Qwen/Qwen2.5-0.5B-Instruct with

LoRA (r=16). Judge: Qwen/Qwen2.5-1.5B-Instruct loaded locally (no API key required).

Training corpus: 16 customer-support scenarios from

`stateset_agents.data.customer_support_bench`; eval: held-out 8 scenarios. Trainer: GSPO

with **safe defaults from §B.1** (`use_reference_model=True`, `beta=0.05`, `num_epochs=1`).

Identical config across all three seeds — the only thing that varies is the RNG state.

Known methodology limitation — in-family judge bias. The trainee is

Qwen2.5-0.5B-Instruct and the judge is Qwen2.5-1.5B-Instruct. The two models share a

training data distribution, tokenizer, and instruction-tuning lineage. In-family LLM-judge

evaluation is known to bias toward outputs that match the judge's own stylistic prior —

the judge "recognizes" its sibling's voice. This does not invalidate the direction of the

result (the rubric agrees independently with +0.072) or its three-seed stability (σ =

0.058) — but the absolute magnitude of `judge_improvement` = +0.079 should be read as

an upper bound on what a cross-family judge would report. The reference notebook supports an optional `CROSS_FAMILY_JUDGE_MODEL` constant (e.g. `meta-llama/Meta-Llama-3-8B-Instruct`) for a sanity check on the same generations; rerun with that set to surface the cross-family number. The v1.1 revision is expected to include the cross-family result inline.

Publication gate. The protocol from [benchmark_results/SCHEMA.md](#) requires judge improvement > 0.03 with three-seed agreement in the positive direction. Both conditions hold ($\Delta = +0.079 > 0.03$; all three seeds positive). The rubric satisfies the gate independently ($\Delta = +0.072 > 0.03$; all three positive).

What this result supports.

1. **§5.1 — GSPO objective.** Group-relative advantages produce reproducible positive transfer across three seeds on a small (16-scenario) corpus.
2. **§10.5 — Reference-model drift.** The framework now has documented evidence across the full headroom spectrum. With no KL anchor on this corpus the policy destabilizes to gibberish ($\Delta = -0.39$, [artifact](#)). With anchor + a near-ceiling base (Qwen3.5-0.8B baseline judge = 0.90) the policy stays within 0.03 of baseline ($\Delta = -0.03$, $\sigma = 0.007$, [artifact](#)). With anchor + a base that has headroom (Qwen2.5-0.5B-Instruct baseline judge = 0.75) the policy improves ($\Delta = +0.08$, this artifact). The KL anchor's "stay close to π_{ref} " property is now first-party-verified across both regimes — *desirable stability* when no headroom, *positive transfer* when headroom exists.
3. **§B.1 — Hyperparameter defaults.** The recommended safe combination (`use_reference_model=True`, `beta=0.05`, all other GSPO defaults paper-recommended) produces the canonical result with zero tuning.

What this result does not support. It does *not* claim that GSPO with these defaults is the best trainer for customer-support tasks at this scale — only that it is *one* configuration that produces reproducible positive transfer. Comparative claims against GRPO / DAPO / VAPO require running the same three-seed protocol against each.

How to reproduce. From a fresh Colab A100 runtime: Runtime → Run all on the reference notebook. The notebook is pinned to a specific framework commit (cell 2) and uses `set_all_seeds()` for canonical determinism on all RNGs.

12. Glossary

Term	Meaning
Advantage	Pinned to source commit. See Appendix C for reproducibility commands. Reward minus baseline; the signal used in the policy gradient. In group-based RL, the baseline is the group mean.
Clip-Higher	Asymmetric PPO clipping with different lower and upper bounds: $\$[1 - \backslash varepsilon_L, 1 + \backslash varepsilon_H]\$$. Introduced by DAPO.

Term	Meaning
Dynamic sampling	A DAPO mechanism that rolls additional samples until every prompt in the batch has at least one correct and one incorrect response, guaranteeing non-zero gradient.
Forward KL	$\mathcal{KL}(\pi_{\theta} \mid \pi_{\text{ref}})$ — the direction used for policy regularization. PyTorch's <code>F.kl_div</code> computes the reverse direction; this framework computes the correct direction manually.
GAE	Generalized Advantage Estimation. A λ -weighted exponential moving average of TD-residuals. VAPO uses λ_{GAE} .
Group-relative advantage	Reward normalized within a group of G samples drawn from the same prompt: $A_i = (r_i - \mu_g) / \sigma_g$.
Importance ratio	$\pi_{\theta}(y) / \pi_{\text{ref}}(y)$
Length-adaptive λ	VAPO's policy GAE λ that grows with sequence length: $\lambda = 1 - 1 / (\alpha L + 1)$, default $\alpha = 0.05$.
LoRA / QLoRA	Low-Rank Adaptation. Trains a small low-rank update to frozen base weights. QLoRA additionally quantizes the base to 4-bit NF4.
Overlong shaping	DAPO's piecewise length penalty: zero in $[0, \text{soft_start}]$, linear in $[\text{soft_start}, \text{max_len}]$, full -1 above max_len .
PEFT	Parameter-Efficient Fine-Tuning (Hugging Face library). Provides LoRA, prefix tuning, and prompt tuning adapters.
Reference model	A frozen snapshot of the policy (typically the SFT initialization) used to anchor the KL regularizer.
Stub backend	Test-only <code>ModelBackend</code> returning deterministic responses without loading any real model. The seam that makes the framework GPU-free for testing.
vLLM	A high-throughput batched generation engine using paged attention. Used as the rollout-phase generator in TRL GRPO and DAPO.

13. Related Work

13.1 Feature comparison matrix

Pinned to source commit. See Appendix G for reproducibility commands. Licensed under the Apache License 2.0 on 2029-01-01.					
Capability	StateSet Agents	TRL (HF)	OpenRLHF	TRLX	NeMo-Aligner
Group-based trainers shipped		GRPO (only)	GRPO	—	GRPO

Capability	StateSet Agents	TRL (HF)	OpenRLHF	TRLX	NeMo-Aligner
	GRPO, GSPO, GEPO, DAPO, VAPO				
Single-turn RLHF (PPO/DPO)	✓ (PPO baseline; DPO via TRL passthrough)	✓ (canonical)	✓	✓	✓
Multi-turn ConversationEnvironment	✓ (first-class, max_turns , scenarios, async step)	✗	partial (Ray actors, no env abstraction)	✗	✗
Tool-using agent (ToolAgent)	✓ (OpenAI-compat function calling, parallel exec)	✗	✗	✗	✗
Offline RL (CQL/IQL/BCQ/BEAR/DT)	✓ (5 algos, blended online/offline schedule)	✗	✗	✓ (one algorithm, ILQL)	✗
Composable reward functions	✓ (CompositeReward , async, graceful degradation)	partial (callable rewards, no composition primitives)	partial	partial	✗
LLM-as-Judge built-in	✓ (LLMJudge , OpenAI/Anthropic/local providers)	✗ (user wires it)	✗	✗	✗
Serving layer	✓ (FastAPI, OpenAI-compat endpoints, Helm chart)	✗	✗	✗	partial (Triton recipe)
Stub-backend testing seam	✓ (StubBackend , 2,438-test suite, no GPU)	✗	✗	✗	✗
Forward-KL analytical (not \$k_1\$/\$k_3\$ approx — see §5.7)	✓ Pinned to source commit. See Appendix C for reproducibility commands. Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).	✗ (\$k_1\$, see ppo_trainer.py:517)	\$k_1\$	\$k_3\$	varies by recipe
Rust-accelerated kernels		✗	✗	✗	✗ (CUDA kernels)

Capability	StateSet Agents	TRL (HF)	OpenRLHF	TRLX	NeMo-Aligner
	✓ (PyO3, GAE + advantage + GSPO ratio)				
License	BUSL-1.1 → Apache 2.0 (2029)	Apache 2.0	Apache 2.0	MIT	Apache 2.0
Distributed training (DeepSpeed / Megatron / Ray)	✓ (Accelerate, DeepSpeed ZeRO 1-3, Ray HPO)	✓ (Accelerate)	✓ (Ray-native)	✓ (Accelerate)	✓ (Megatron-native)
vLLM rollout integration	✓ (TRL GRPO + DAPO paths)	✓ (recent)	✓	✗	partial
Production multi-turn conversation memory	✓ (5-tier; Redis/SQLite backends)	✗	✗	✗	✗

✓ = first-class supported. partial = supported through user wiring. ✗ = not present.

13.2 Distinguishing position

The matrix is unsentimental about where StateSet Agents adds nothing: distributed training, single-turn RLHF, basic vLLM integration are all well-served elsewhere. The framework's *distinctive* position lives in the column-pairs:

- Five group-based trainers + multi-turn environments + serving layer** in one library. No other framework in the matrix treats all three as first-class concerns in a single deployable package — OpenRLHF has Ray actors and a serving recipe, NeMo-Aligner has a Triton recipe, TRL has GRPO; each treats one or two of these as primary and the others as user-wired.
- Composable async rewards + LLM-judge built-in + stub-backend testing.** Lets you iterate on reward design at full speed without needing a GPU.
- Tool-using agents as first-class RL targets**, with the tool-call semantics pinned down (§3.1) instead of left to user interpretation.

If your problem is single-turn preference data on Llama, use DPO in TRL. If your problem is multi-turn dialogue with tools, where the reward function is itself an experimental artifact, the matrix above explains why this framework exists.

Pinned to source commit. See Appendix C for reproducibility commands.

Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

14. Conclusion

StateSet Agents packages recent advances in group-based policy optimization into a single, deployable framework. By treating multi-turn conversations as the unit of

optimization, by making rewards composable and async, by separating algorithm from environment from agent from backend, and by investing in both training and serving, the framework lets practitioners go from a Hugging Face checkpoint to a trained, served, monitored conversational agent without crossing a tool boundary. The five trainers — GRPO, GSPO, GEPO, DAPO, VAPO — cover a spectrum from baseline simplicity to the strongest reported reasoning results, and the stub-backend testing pattern keeps the iteration loop tight even without GPU access.

The five design principles that anchor the framework are stated in **§1.0 (Design Philosophy)** at the top of this document. They are the same five principles whether you read them as a preface or as a closing summary; we put them up front because they constrain everything between.

References

1. Zheng et al. *Group Sequence Policy Optimization*. arXiv:2507.18071, 2025.
2. Yu et al. *DAPO: An Open-Source LLM Reinforcement Learning System at Scale*. arXiv:2503.14476, 2025.
3. Yue et al. *VAPO: Efficient and Reliable Reinforcement Learning for Advanced Reasoning Tasks*. arXiv:2504.05118, 2025.
4. *Group Expectation Policy Optimization*. arXiv:2508.17850, 2025.
5. Shao et al. *DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models*. arXiv:2402.03300, 2024.
6. Schulman et al. *Proximal Policy Optimization Algorithms*. arXiv:1707.06347, 2017.
7. Schulman et al. *High-Dimensional Continuous Control Using Generalized Advantage Estimation*. arXiv:1506.02438, 2015. (GAE)
8. Ouyang et al. *Training Language Models to Follow Instructions with Human Feedback*. arXiv:2203.02155, 2022.
9. Bai et al. *Constitutional AI: Harmlessness from AI Feedback*. arXiv:2212.08073, 2022.
10. Lee et al. *RLAIF: Scaling Reinforcement Learning from Human Feedback with AI Feedback*. arXiv:2309.00267, 2023.
11. Rafailov et al. *Direct Preference Optimization*. arXiv:2305.18290, 2023. (For contrast with the on-policy methods used here.)
12. Hu et al. *LoRA: Low-Rank Adaptation of Large Language Models*. arXiv:2106.09685, 2021.
13. Dettmers et al. *QLoRA: Efficient Finetuning of Quantized LLMs*. arXiv:2305.14314, 2023. Prior to April 2024, this Appendix C for Quantizability to Meads. Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).
14. Kwon et al. *Efficient Memory Management for Large Language Model Serving with PagedAttention*. SOSP 2023. (vLLM)

15. Kumar et al. *Conservative Q-Learning for Offline Reinforcement Learning*. NeurIPS 2020. (CQL)
16. Kostrikov et al. *Offline Reinforcement Learning with Implicit Q-Learning*. arXiv: 2110.06169, 2021. (IQL)
17. Chen et al. *Decision Transformer: Reinforcement Learning via Sequence Modeling*. NeurIPS 2021.
18. Kirkpatrick et al. *Overcoming Catastrophic Forgetting in Neural Networks*. PNAS 2017. (EWC)
19. Li & Hoiem. *Learning without Forgetting*. ECCV 2016. (LwF)
20. Sutton, R. *The Bitter Lesson*. Personal essay, 2019. <http://www.incompleteideas.net/IncIdeas/BitterLesson.html>

Appendix A: File Map

Component	Path
Base trainer	<code>stateset_agents/training/base_trainer.py</code>
TRL GRPO	<code>stateset_agents/training/trl_grpo_trainer.py</code>
GSPO	<code>stateset_agents/training/gspo_trainer.py</code>
GEPO	<code>stateset_agents/training/gepo_trainer.py</code>
DAPO	<code>stateset_agents/training/dapo_trainer.py</code>
VAPO	<code>stateset_agents/training/vapo_trainer.py</code>
Agent base	<code>stateset_agents/core/agent.py</code>
MultiTurnAgent	<code>stateset_agents/core/multiturn_agent.py</code>
ToolAgent	<code>stateset_agents/core/tool_agent.py</code>
AgentConfig	<code>stateset_agents/core/agent_config.py</code>
Backends	<code>stateset_agents/core/agent_backends.py</code>
Environment base	<code>stateset_agents/core/environment_base.py</code>
ConversationEnvironment	<code>stateset_agents/core/conversation_environment.py</code>
Reward base	<code>stateset_agents/core/reward_base.py</code>
Basic / domain rewards	<code>stateset_agents/core/basic_rewards.py</code> , <code>domain_rewards.py</code>
Multi-objective reward	<code>stateset_agents/rewards/multi_objective_reward.py</code>
Memory	<code>stateset_agents/core/memory.py</code>

Pinned to source commit. See Appendix C for reproducibility commands.
 Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

Component	Path
Neural reward trainer	stateset_agents/training/neural_reward_trainer.py
RLAIF trainer	stateset_agents/training/rlaif_trainer.py
Offline GRPO trainer	stateset_agents/training/offline_grpo_trainer.py
Continual learning	stateset_agents/training/continual_learning.py
Sim-to-real	stateset_agents/training/sim_to_real.py
Domain randomization	stateset_agents/training/domain_randomization.py
Auto-research	stateset_agents/training/auto_research/
Rust acceleration	rust_core/ (crate stateset-rl-core)
FastAPI app	stateset_agents/api/main.py
GRPO metrics	stateset_agents/api/grpo/metrics.py
API middleware	stateset_agents/api/middleware.py
Dashboard	dashboard/ (React 19 + Vite)
Exception taxonomy	stateset_agents/exceptions.py
Helm chart	deployment/helm/
K8s manifests	deployment/kubernetes/

Appendix B: Hyperparameter Reference

Defaults are taken verbatim from the corresponding config dataclasses in `stateset_agents/training/`. Values that differ between trainers are flagged in the per-trainer tables.

B.1 Shared Defaults (BaseTrainerConfig)

Field	Default	Notes
<code>learning_rate</code>	<code>1e-5</code>	AdamW base LR
<code>adam_beta1</code>	<code>0.9</code>	AdamW β_1
<code>adam_beta2</code>	<code>0.999</code>	AdamW β_2
<code>weight_decay</code>	<code>0.01</code>	L2 regularization
<code>max_grad_norm</code>	<code>1.0</code>	Gradient clipping threshold

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

Field	Default	Notes
warmup_ratio	0.1	Cosine warmup fraction
num_epochs	3	Training epochs
num_episodes	100	Total episodes generated
bf16	True	bfloat16 mixed precision (H100/A100 default)
fp16	False	float16 fallback for older GPUs
gradient_checkpointing	True	Memory-saving recomputation
use_4bit	False	QLoRA NF4 quantization
use_8bit	False	8-bit quantization
use_lora	True	LoRA adapters on by default
lora_r	16	LoRA rank
lora_alpha	32	LoRA scaling
lora_dropout	0.05	LoRA dropout
lora_target_modules	None	Defaults to model-appropriate projection layers
use_vllm	False	vLLM rollout acceleration
vllm_gpu_memory_utilization	0.85	vLLM memory cap
vllm_tensor_parallel_size	1	TP shards
vllm_enable_prefix_caching	True	vLLM prompt-prefix cache
beta	0.0	KL penalty coefficient (off by default). See warning below.
use_reference_model	False	Frozen π . See warning below.
max_prompt_length	256	Prompt token cap
max_completion_length	512	Response token cap
temperature	0.7	Sampling temperature
top_p	0.9	Nucleus sampling
logging_steps	10	Metric logging interval
save_steps	500	Checkpoint interval
eval_steps	100	Evaluation interval
report_to	"wandb"	Logging backend

Pinned to source commit. See Appendix C for reproducibility commands.
 Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

Warning on `beta=0.0 / use_reference_model=False`. These defaults are correct for production-scale training where the rollout count is large enough that group-relative advantages alone handle policy drift. On small corpora (rule of thumb: fewer than ~100 training queries) this combination has been observed to destabilize the policy — the trained model emits incoherent token soup while still scoring nonzero on rule-based rewards (see §10.5 for the worked example). A safe default for small-corpus runs is `use_reference_model=True, beta=0.05`. `train_with_gspo` emits a runtime warning when the unsafe combination is detected on a corpus below the threshold.

B.2 TRL GRPO

Field	Default
<code>num_generations</code>	4
<code>num_iterations</code>	1
<code>mini_batch_size</code>	1
<code>num_outer_iterations</code>	1
<code>generations_per_iteration</code>	100
<code>beta</code>	0.0

B.3 GSPO

Field	Default	Notes
<code>num_generations</code>	4	Group size G
<code>clip_range_left</code>	3e-4	ϵ_L for sequence-level ratio (note: very tight)
<code>clip_range_right</code>	4e-4	ϵ_H for sequence-level ratio
<code>num_iterations</code>	1	No inner PPO epochs
<code>mini_batch_size</code>	1	—
<code>use_gspo_token</code>	False	Token-level variant (off by default; sequence-level wins)

Note on the tight clip range. Because GSPO ratios are length-normalized (already exp of a small per-token quantity), the effective clip needs to be much tighter than token-level PPO's 0.2. The defaults 3e-4 / 4e-4 are inherited from the paper. If you see no exploration, widen these first.

B.4 GEPO

Field	Default	Notes
<code>group_size</code>	8	Larger than GRPO/GSPO; group expectation needs more samples
<code>clip_eps</code>	0.2	Standard PPO clip on the GEPO coefficient
<code>learning_rate</code>	1e-6	Paper-recommended; lower than base default
<code>warmup_ratio</code>	0.03	Short warmup
<code>per_device_train_batch_size</code>	8	—
<code>gradient_accumulation_steps</code>	8	Effective batch = 64 per device
<code>use_group_baseline</code>	True	Within-group advantage normalization

B.5 DAPO

Field	Default	Notes
<code>group_size</code>	16	Larger than GRPO/GSPO — dynamic sampling needs candidates
<code>prompt_batch_size</code>	512	Unique prompts per outer batch
<code>mini_batch_size</code>	512	—
<code>num_gradient_updates</code>	16	Multiple inner updates per rollout
<code>clip_eps_low</code>	0.2	ϵ_L
<code>clip_eps_high</code>	0.28	ϵ_H — the Clip-Higher asymmetry
<code>use_dynamic_sampling</code>	True	Filter out 0%- and 100%-accuracy groups
<code>min_accuracy_threshold</code>	0.0	Strict < filter
<code>max_accuracy_threshold</code>	1.0	Strict < filter
<code>dynamic_sampling_buffer_size</code>	1024	Rolling buffer capacity
<code>use_overlong_shaping</code>	True	Piecewise length penalty
<code>max_generation_length</code>	20480	Hard length cap
<code>overlong_cache_length</code>	4096	Linear-penalty region width
<code>overlong_penalty</code>	-1.0	Penalty at and beyond <code>max_length</code>
<code>use_token_level_loss</code>	True	Divide by total response tokens
<code>learning_rate</code>	1e-6	Per paper

Pinned to source commit. See Appendix C for reproducibility commands.
 Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03)

Field	Default	Notes
<code>lr_scheduler_type</code>	"constant"	Per paper
<code>temperature</code>	1.0	Higher than base default
<code>top_p</code>	0.7	—
<code>eval_repeats</code>	32	Stochastic-eval averaging

B.6 VAPO

Field	Default	Notes
<code>group_size</code>	16	Same as DAPO
<code>num_prompts_per_batch</code>	512	Outer batch
<code>mini_batch_size</code>	512	—
<code>clip_eps_low</code>	0.2	Same Clip-Higher as DAPO
<code>clip_eps_high</code>	0.28	—
<code>use_token_level_loss</code>	True	—
<code>value_hidden_size</code>	1024	ValueHead MLP width
<code>value_num_layers</code>	2	ValueHead depth
<code>value_warmup_steps</code>	50	Critic-only warmup before joint training
<code>lambda_critic</code>	1.0	Unbiased MC for the critic
<code>lambda_policy_alpha</code>	0.05	Length-adaptive λ slope
<code>value_loss_coef</code>	0.5	Weight on $\mathcal{V}$
<code>entropy_coef</code>	0.0	Entropy bonus (off by default)
<code>use_positive_lm_loss</code>	True	NLL on correct samples
<code>positive_lm_weight</code>	0.1	Weight on positive-LM term
<code>actor_learning_rate</code>	1e-6	Per paper
<code>critic_learning_rate</code>	2e-6	Slightly higher than actor

B.7 Offline GRPO

Pinned to source commit. See Appendix C for reproducibility commands.

Field	Default	Notes
<code>offline_algorithm</code>	"iql"	One of: <code>cql</code> , <code>iql</code> , <code>bcq</code> , <code>bear</code> , <code>dt</code>
<code>value_hidden_size</code>	256	—

Licensed under the Apache License 2.0 (transitioning to Apache 2.0 on 2029-09-03).

Field	Default	Notes
value_num_layers	3	—
value_activation	"relu"	One of: relu , tanh , gelu
offline_weight	0.5	Weight on offline value estimate
online_weight	0.5	Weight on online group baseline
warmup_offline_steps	1000	Pure-offline phase length
blend_schedule	"linear"	One of: linear , exponential , constant
num_generations	8	Online phase group size
clip_ratio	0.2	PPO clip on online GRPO updates
baseline_type	"hybrid"	One of: offline , online , hybrid
pretrain_value_epochs	10	Value-only pretraining epochs
value_learning_rate	3e-4	Separate from policy LR
embedding_model	"all-MiniLM-L6-v2"	Sentence-transformer for state/action embeddings
state_dim	384	MiniLM output dim
action_dim	384	—

B.8 Sizing Cheatsheet

A rough memory budget for a 7B-parameter base model with LoRA $r=16$:

Component	Memory
Base model (bf16)	~14 GB
Base model (4-bit NF4)	~5 GB
LoRA adapters ($r=16$)	~150 MB
Reference model (bf16, frozen)	~14 GB
ValueHead (VAPO only)	~10 MB
Optimizer state (AdamW, LoRA only)	~300 MB
Activations (gradient checkpointing)	~2-6 GB depending on context length

Pinned to source commit. See Appendix C for reproducibility commands.

This means a 7B model with QLoRA + reference model fits in ~25 GB during training — single 40 GB A100 territory. A 30B model with the same configuration needs ~60 GB — H100 or multi-GPU. VAPO adds the ValueHead but no significant memory burden compared to keeping a reference model.

Appendix C: Reproducibility Commands

Every claim in this whitepaper that names a file path, line number, default value, or test count is verifiable from a checkout of commit `4744c76`. The commands below are the canonical way to verify each class of claim.

C.1 Checkout

```
git clone https://github.com/stateset/stateset-agents
cd stateset-agents
git checkout 4744c76
```

C.2 Verify implementation citations

```
# §5.2 – GSP0 sequence-importance-ratio computation
sed -n '390,419p' stateset_agents/training/gspo_trainer.py

# §5.3 – GEP0 group-expectation denominator
sed -n '301,337p' stateset_agents/training/gepo_trainer.py

# §5.4 – DAP0 dynamic-sampling filter
sed -n '276,278p' stateset_agents/training/dapo_trainer.py

# §5.5 – VAP0 ValueHead + LengthAdaptiveGAE + warmup
sed -n '177,349p' stateset_agents/training/vapo_trainer.py
sed -n '575,680p' stateset_agents/training/vapo_trainer.py

# §5.7 – Forward-KL computation
sed -n '600,631p' stateset_agents/training/base_trainer.py
```

C.3 Verify defaults (Appendix B)

```
# Every default named in Appendix B can be located via:
grep -rn "default=\|: int = \|: float = \|: bool = " stateset_agents/training/ | less

# Specific examples:
grep -n "clip_range_left\|clip_range_right" stateset_agents/training/gspo_config.py
grep -n "clip_eps_low\|clip_eps_high\|group_size" stateset_agents/training/dapo_config.py
grep -n "value_warmup_steps\|lambda_policy_alpha" stateset_agents/training/vapo_config.py
```

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

C.4 Verify test count and coverage methodology

```
# Test count + green status. Confirms both the count AND that the suite passes –
# what readers actually care about for an audit trail.
pytest tests/unit -q --no-header 2>&1 | tail -2

# Coverage (claimed: ~49% overall on in-process paths)
pytest --cov=stateset_agents --cov-report=term tests/unit tests/integration 2>&1 | tail -5
```

C.5 Run the benchmark methodology (§7.5)

```
pip install -e ".[training,vllm]"

# End-to-end latency/throughput
python benchmarks/performance_benchmarks.py --latency --throughput --report html

# Head-to-head GRPO vs GSP0 vs DAPO
python benchmarks/algorithm_comparison.py

# vLLM vs HF generation
python benchmarks/real_performance_benchmarks.py
```

Output lands in `benchmark_results/` as JSON + (optionally) HTML. **Record your hardware, CUDA version, driver, model checkpoint, and random seed alongside the numbers** — these dominate the reported throughput more than any framework-level decision.

C.6 Verify Helm chart and K8s manifests

```
helm lint deployment/helm/
helm template deployment/helm/ -f deployment/helm/values-a100.yaml > /tmp/a100.yaml
helm template deployment/helm/ -f deployment/helm/values-h100.yaml > /tmp/h100.yaml
```

C.7 Audit trail summary

Claim	Where to verify
Trainer line counts	<code>wc -l stateset_agents/training/*_trainer.py</code>
Default hyperparameters	<code>*_config.py</code> dataclass fields
File paths in Appendix A	<code>find stateset_agents/ -name '*.py' \ grep <component></code>
Exception tuple definitions	<code>cat stateset_agents/exceptions.py</code>
Rust functions exposed to Python	<code>cat rust_core/src/lib.rs</code>
Prometheus metric names	<code>grep -rn "Counter\ Histogram\ Gauge" stateset_agents/api/</code>

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).

Claim	Where to verify
Helm values overlays	<code>ls deployment/helm/values-*.yaml</code>
K8s manifests	<code>ls deployment/kubernetes/</code>

If any of these commands return output that disagrees with what this whitepaper says, **the code wins** — please open an issue referencing this whitepaper version and commit hash so we can correct the document.

Copyright © 2026 StateSet. Licensed under the Business Source License 1.1, converting to Apache 2.0 on 2029-09-03.

1. The BUSL period reserves hosted-derivative offering for StateSet; individual users, academic research, and customers building on top of the framework have always-permitted use under the [Additional Use Grant](#). The Apache 2.0 conversion on 2029-09-03 is a binding contractual term, not a marketing statement. ↩

Pinned to source commit. See Appendix C for reproducibility commands.
Licensed BUSL-1.1 (transitioning to Apache 2.0 on 2029-09-03).